# An Introduction to

Graphics Programming

with

OpenGL™

Tutorial and Reference Manual

Toby Howard

Department of Computer Science
University of Manchester

Contents

5   2D and 3D graphics   **21**

6   Animated graphics   **27**

PART II   OpenGL Reference Manual   **33**

7   Graphics primitives   **35**

8   Modelling using transformations   **45**

Chapter 1

About this manual

This manual is in two parts: the first (Chapters 2 to 6) is a hands-on **Tutorial**, which uses a series of example programs to illustrate some of the main features of OpenGL. The second part (Chapter 7 onwards) is a **Reference Manual**, which describes some OpenGL functions in detail.

1.1   How to read this manual

If you're a newcomer to OpenGL, we recommend that you first read the tutorial chapters, in order, and experiment with the example programs on-line. These chapters introduce the basic concepts of OpenGL, and cover the details of how to compile and run OpenGL C programs using our local GNU/Linux installation.

The reference chapters are intended to support the lecture material and the laboratory programming exercises.

1.2   Join the Bug Club

In the highly unlikely event that you find a bug in this manual, please email us the details. Successful correspondents will receive honorary membership of the Bug Club. Send bug reports to `opengl@cs.man.ac.uk`.

1.3   Acknowledgements

It's a pleasure to thank Alan Murta and Julien Cartigny for helping with parts of this manual. And thank you to all the people who have made their excellent GNU/Linux software freely available: Mesa (which includes GLU) was written by Brian Paul (http://www.mesa3d.org). GLUT was written by Mark J. Kilgard, who kindly provided additional help.

Part I

OpenGL Tutorial

Chapter 2

Introduction to OpenGL

In recent years OpenGL has become a worldwide standard for 3D computer graphics programming. It's very widely used: in industry, in research laboratories, in computer games – and for teaching computer graphics.

OpenGL is a powerful, professional-level system, and it would take a manual much thicker than this one to describe all its facilities completely. We have selected a **subset** of OpenGL – a portion of OpenGL's functionality which is relevant to the CS2072 Interactive Graphics course, and sufficient to support its programming labs.

2.1 What is OpenGL?

OpenGL has it origins in the earlier GL ("Graphics Library") system which was invented by Silicon Graphics Inc. as the means for programming their high-performance specialised graphics workstations. As time went on, people became interested in porting GL to other kinds of machine, and in 1992 a variation of GL – called OpenGL – was announced. Unlike GL, OpenGL was specifically designed to be **platform-independent**, so it would work across a whole range of computer hardware – not just Silicon Graphics machines. The combination of OpenGL's power and portability led to its rapid acceptance as a **standard** for computer graphics programming.

OpenGL itself isn't a programming language, or a software library. It's the **specification** of an Application Programming Interface (API) for computer graphics programming. In other words, OpenGL defines a set of functions for doing computer graphics.

What you actually use to do your graphics is an **implementation** of OpenGL. We use a free software system called **Mesa**, which we'll describe in Section 2.3.

2.2 A whirlwind tour of OpenGL

What exactly can OpenGL do? Here are some of its main features:

- It provides 3D geometric objects, such as lines, polygons, triangle meshes, spheres, cubes, quadric surfaces, NURBS curves and surfaces;

5
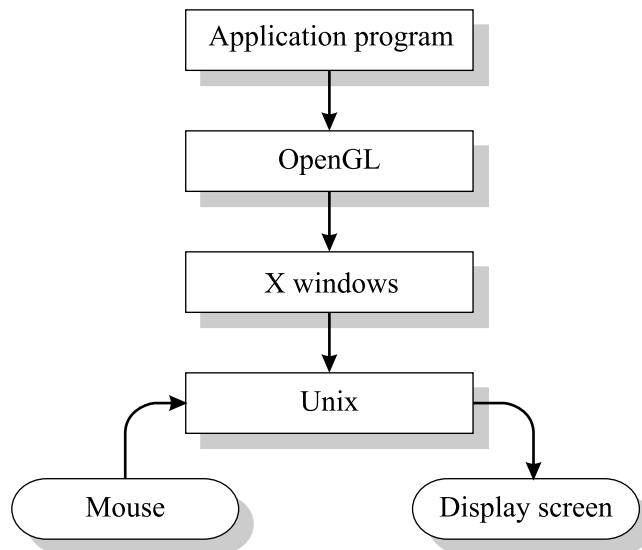
Figure 2.1: Where OpenGL fits in – a high-level view.

- It provides 3D modelling transformations, and viewing functions to create views of 3D scenes using the idea of a **virtual camera**;

- It supports high-quality rendering of scenes, including hidden-surface removal, multiple light sources, material types, transparency, textures, blending, fog;

- It provides display lists for creating graphics caches and hierarchical models. It also supports the interactive "picking" of objects;

- It supports the manipulation of images as pixels, enabling frame-buffer effects such as anti-aliasing, motion blur, depth of field and soft shadows.

Figure 2.1 shows the relationship between an application and OpenGL in our local GNU/Linux environment. An application programmer sees OpenGL as a single library providing a set of functions for graphical input and output. In fact, it's slightly more complicated than that.

2.2.1   The support libraries: GLU and GLUT

A key feature of the design of OpenGL is the separation of **interaction** (input and windowing functions) from **rendering**. OpenGL itself is concerned only with graphics rendering. You can always identify an OpenGL function: all OpenGL function names start with **"gl"**.

Over time, two **utility libraries** have been developed which greatly extend the low-level (but very efficient) functionality of OpenGL. The first is the "OpenGL Utility Library", or **GLU**. The second is the "OpenGL Utility Toolkit", or **GLUT**:

- **GLU** provides functions for drawing more complex primitives than those of OpenGL, such as curves and surfaces, and also functions to help specify 3D views of scenes. All GLU function names start with **"glu"**.
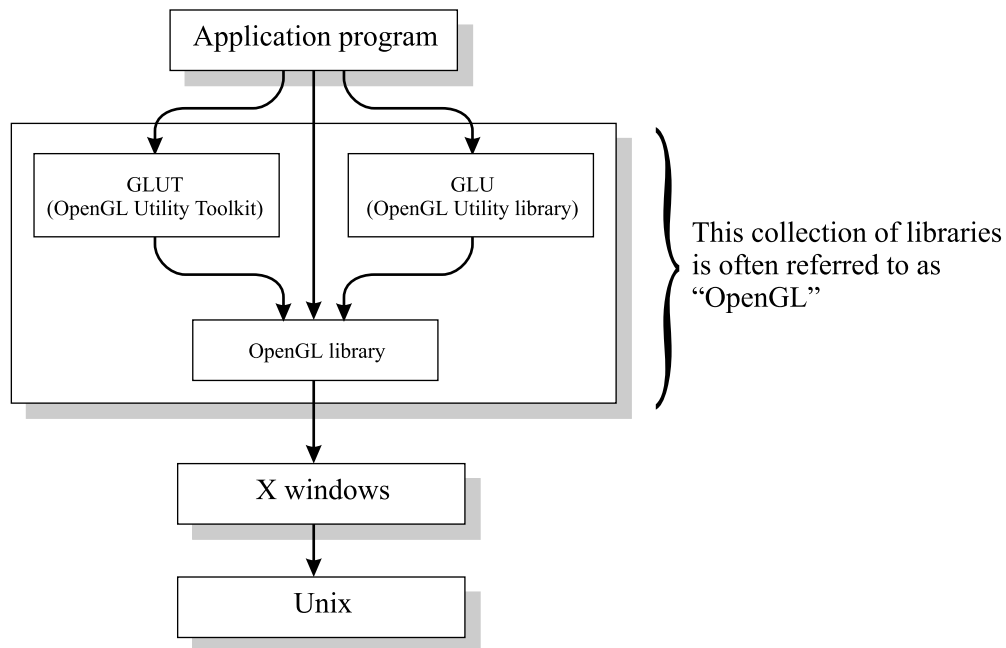
Figure 2.2: What is commonly called "OpenGL" is actually a set of three libraries: **OpenGL** itself, and the supporting libraries **GLU** and **GLUT**.

- **GLUT** provides the facilities for interaction that OpenGL lacks. It provides functions for managing windows on the display screen, and handling input events from the mouse and keyboard. It provides some rudimentary tools for creating Graphical User Interfaces (GUIs). It also includes functions for conveniently drawing 3D objects like the platonic solids, and a teapot. All GLUT function names start with **"glut"**.

Figure 2.2 shows the relationships between OpenGL, GLU, and GLUT. As you can see, it's helpful to think of "layers" of software, where each layer calls upon the facilities of software in a lower layer.

However, somewhat confusingly, when most people say "**OpenGL**", what they really mean is "**OpenGL** plus **GLU** plus **GLUT**". It's a slightly lazy terminology, but we'll use it too.

2.3   What is Mesa?

Mesa is a C implementation of a graphics system that looks **extremely similar** to the official OpenGL specification. (We can't actually say "Mesa **is** an implementation of OpenGL" for legal reasons. But, for all intents and purposes, it is really.)

Whereas OpenGL is intended to run on machines which have graphics support in hardware, Mesa doesn't require the presence of any special 3D graphics acceleration hardware – although it can certainly take advantage of it if it's there. Of course, the performance of the graphics will be better with hardware acceleration, but it's still remarkably good without it on a reasonably fast PC.

Mesa is free "open software". Hurrah!

2.4   Using OpenGL away from the Department

Mesa has been ported to many different platforms, including GNU/Linux, SunOS, DOS, Windows, and OS/2. In the Department, however, we currently support Mesa **only on GNU/Linux**.

If you wish to run Mesa on GNU/Linux away from the Department, refer to our local OpenGL Web pages (see Section 2.5.1), which explain where to get the software, and give some installation guidelines.

For any other platform – specifically **Windows** – see the next section for pointers to resources, but otherwise you're on your own.

2.5   Resources and further reading

Here are some useful resources, and suggestions for further reading, should you wish to find out more.

2.5.1   On-line resources

- Our local OpenGL Web pages: **www.cs.man.ac.uk/applhax/OpenGL**. Check here for up-to-date details of the local installation.

- Local example programs: we have a number on-line, in **/opt/info/courses/OpenGL/examples**.

- The local newsgroup for graphics queries and chat: **man.cs.graphics**. Post your OpenGL programming queries here, and help others with theirs.

- The USENET OpenGL newsgroup: **comp.graphics.api.opengl**. This can be a great source of help and information, for newcomers and experts alike. However, note that it is **highly inadvisable** to post pages of source code saying "my program doesn't work". As with all newsgroups, lurk for a while and get a feel of the etiquette before posting.

- The official home of OpenGL on the Web: **www.opengl.org**. Lots of pointers to on-line information, tutorials, example programs, and downloadable software.

2.5.2   Books

- **Interactive Computer Graphics: A Top-Down Approach with OpenGL** by Edward Angel. Addison-Wesley, ISBN 0-201-85571-2. General introduction to computer graphics for people new to the subject. This is the recommended textbook for the CS2072 course.

- **OpenGL Programming Guide, Third Edition: The Official Guide to Learning OpenGL, Version 1.2** by Mason Woo, Jackie Neider, Tom Davis and Dave Shreiner. Addison-Wesley, ISBN 0-201-60458-2. Also known as "The Red Book", provides complete coverage of OpenGL from simple to advanced, with many code examples. Assumes familiarity with C, some maths, geometry. The coverage of this book far exceeds the material taught in CS2072. Beware: the book contains a few errors, so check the errata list at **http://www.woo.com/errata.html**.

- **Programming OpenGL for the X Window System** by Mark J. Kilgard. Addison-Wesley, ISBN: 0-201-48359-9.

- **OpenGL Superbible: The Complete Guide to OpenGL Programming for Windows NT and Windows 95** by Richard S. Wright Jr. and Michael Sweet, Waite Group Press, ISBN 1-57169-073-5. Describes how to use OpenGL on Windows 95/98/NT, with things like Visual C++, MFC, Borland 5.0 and OWL. Beware: contains errors in its description of OpenGL, and American humour. Note: this book is **not relevant** to using OpenGL on GNU/Linux.

2.5.3   Technical documentation

- **The OpenGL Graphics System: A Specification (Version 1.2)** by Mark Segal and Kurt Akeley. On-line at **trant.sgi.com/opengl/docs/Specs/glspec1.1/glspec.html**, printed copies available on loan from the Student Resource Centre. This manual is the detailed (and highly technical) specification of OpenGL – for advanced readers only.

- **The OpenGL Utility Toolkit (GLUT) Programming Interface (API Version 3)** by Mark J. Kilgard. On-line at **reality.sgi.com/opengl/spec3/spec3.html**, printed copies available on loan from the Student Resource Centre. This manual is the detailed technical specification of GLUT.

2.6   About the notation used in this manual

Experienced C programmers might wish to skip this section.

In this manual, when we introduce a new OpenGL function, we'll give its definition, followed immediately by a description of what it does.

To take an example at random, here's the definition of the GLUT function which draws a sphere, which you'll meet on page 42:

```
void glutWireSphere ( GLdouble radius,
                      GLint slices,
                      GLint stacks );
```

What this notation means is the following:

- The name of the function is **glutWireSphere**();

- The result type of the function is `void`;

- The function has three arguments:

    - `radius`, of type `GLdouble`
    - `slices`, of type `GLint`
    - `stacks`, of type `GLint`

To actually **use** this function in your program, you would do something like this:

```
GLdouble rad= 1.0;
GLint sl= 15;
GLint st= 20;

glutWireSphere (rad, sl, st);
```

Or, you could set the arguments directly, without declaring variables:

```
glutWireSphere (1.0, 15, 20);
```

Note that OpenGL defines its own names for data types, all of which begin with `GL`. Examples are: `GLdouble`, `GLint`, `GLfloat`. The reason it's done like this is to make the specification of OpenGL language-independent. In most cases, it'll be obvious what the data type means – `GLint`, for example, is GL's name for an integer, or an `int` in C. Where it isn't obvious, we'll tell you.

To continue with the example of **glutWireSphere()**, this is how we'd write its description:

> **glutWireSphere()** draws a sphere, of radius `radius`, centred on $(0, 0, 0)$ in object co-ordinates. `slices` is the number of subdivisions around the Z axis (like lines of longitude); `stacks` is the number of subdivisions along the Z axis (like lines of latitude). Solid version: **glutSolidSphere()**.

2.7   What next?

Now onto Chapter 3, which explains how to compile OpenGL programs using our local installation.

Chapter 3

Getting started with OpenGL

This chapter explains how to compile and link C programs with OpenGL using our local installation. There are two different ways to do this:

- Using the command `cogl` – **this is handy for compiling single standalone OpenGL programs, and is the recommended way for compiling programs in the CS2072 lab**; (`cogl` is a Perl script and lives in `/usr/local/bin`).

- Using a makefile – this is a more flexible approach, necessary for larger projects which use more than one source file. Use of a makefile is **not recommended** for the CS2072 lab. See Appendix B for a sample makefile.

3.1   Compiling using `cogl`

`cogl` is a command we've written locally to make compiling single programs with OpenGL as simple as possible. (The Perl source code of cogl is listed in Appendix A).

We'll use the example program `thegears.c` to illustrate the use of `cogl`.

First, make sure you are running X Windows. Then, select an appropriate directory to work in, and take your own private copy of the program `thegears.c`, as follows (the string `punter$` stands for whatever command prompt your shell window displays):

```
punter$ cp /opt/info/courses/OpenGL/examples/thegears.c .
```

(Don't forget that **dot (.)** as the second argument to `cp`.)

You compile and link the program as follows:

```
punter$ cogl thegears.c
```

This will produce an executable program called `thegears`, which you run as follows:

```
punter$ thegears
```

You should see a square OpenGL window appear on your display, with something interesting happening within it. Move your mouse into the OpenGL window, and press `h` on the keyboard to bring up the help screen. Experiment with the program as it suggests.

3.2 Try some other examples

There are a number of other example programs in `/opt/info/courses/OpenGL/examples/`, which we'd encourage you to copy, compile and play with. Here are some we recommend.

- `tori`: some doughnuts. Move the mouse slowly;

- `teapots`: draws our teapot collection;

- `morph3d`: might remind you of a certain screensaver;

- `reflect`: reflective texture mapping. Try the arrow keys;

- `pointblast`: a simple particle system. Try the mouse buttons;

- `star`: moving starfield. Hit `t` to warp;

- `lorenz`: chaos. Have aspirins handy.

3.3 Yet more examples

Here are some other examples to try, again in `/opt/info/courses/OpenGL/examples/`. These are part of the `xscreensaver` collection, and are already compiled for you, so just `cd` to that directory, and run the programs. You'll have to type `control-c` in your shell to stop them running:

- `moebius`: ants crawl inexplicably around a moebius strip;

- `sproingies`: multi-coloured bouncy things tumble down an infinite staircase, and occasionally explode;

- `superquadrics`: 3D shapes morph into each other, based on the "superquadric" objects developed by American graphics researcher Alan Barr;

- `cage`: be amazed as OpenGL draws an impossible object.

3.4 What next?

Now onto Chapter 4, which introduces the structure of an OpenGL program.

Chapter 4

Beginning OpenGL programming

In this and the next two chapters, we introduce the basic ideas of OpenGL in a tutorial fashion, using a series of example programs.

4.1 The OpenGL model

Figure 4.1 shows the relationships between an application program, the graphics system, input and output devices, and the user.

The **application program** has its own internal **model** of what it's doing – its own interpretation of what the graphics it's manipulating actually **means**. It draws the graphics using the facilities of the **graphics system** – in our case, OpenGL. The user views the graphics, and uses **input devices**, such as a mouse, to **interact**. Information about the user's interactions are sent back to the application, which decides what action to take. Typically, it will make changes to its internal model, which will cause the graphics to be updated, and so another **loop** in the interaction cycle begins.

4.2 Platform- and device-independence

As we saw in Chapter 2, OpenGL is designed to be platform-independent and device-independent, so it isn't concerned with the exact makes and models of graphics display and interaction hardware it uses. Instead, OpenGL functions refer to **windows** and **events**:

- An OpenGL **window** is a rectangular area on a physical display screen into which OpenGL draws graphics. Usually, an OpenGL window corresponds exactly to a window managed by the "window manager", such as X. (It's also possible to have multiple OpenGL windows simultaneously active on a single display – see Section 9.4.)

- An OpenGL **event** occurs when the user operates an input device. In order to respond to the input event, the application must provide a C function – known as a **callback function** – to handle the event; OpenGL automatically calls the application's function, passing it the event data.
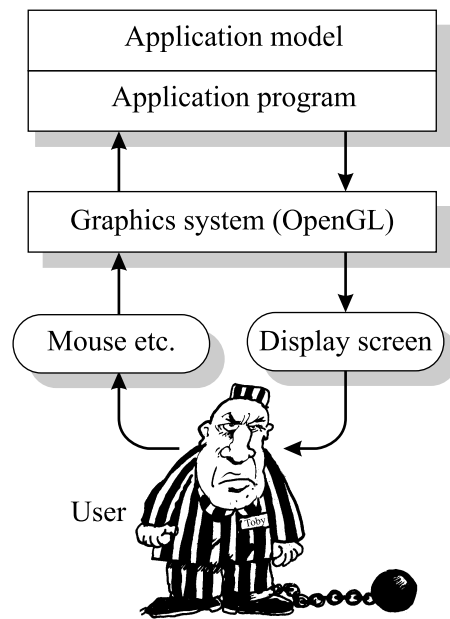
Figure 4.1: The graphical interaction loop.

In fact, OpenGL doesn't draw its graphics directly to the window. It actually draws into a data structure (an array of pixels) inside OpenGL called the **frame-buffer**, often just called the **buffer**. Periodically, OpenGL copies the pixels in the frame buffer into the window. More on this in Section 6.2.

4.3   Example 1: a bare-bones program

We'll begin with the simplest possible OpenGL program. It's `ex1.c` in the `examples` directory. Take a copy of this program, and compile it with `cogl`:

```
punter$ cp /opt/info/courses/OpenGL/examples/ex1.c .
punter$ cogl ex1.c
```

When you run `ex1`, you should see an OpenGL window appear. To stop the program running, place your mouse inside the shell window from which you ran the program, and hit `control-c`.

Here's the code for `ex1.c`:

```
/* ex1.c */
#include <GL/glut.h>

void display (void) {
/* Called when OpenGL needs to update the display */
  glClear (GL_COLOR_BUFFER_BIT); /* Clear the window */
  glFlush();                     /* Force update of screen */
}
```

```
int main (int argc, char **argv) {
  glutInit (&argc, argv);     /* Initialise OpenGL */
  glutCreateWindow ("ex1");  /* Create the window */
  glutDisplayFunc (display); /* Register the "display" function */
  glutMainLoop ();            /* Enter the OpenGL main loop */
  return 0;
}
/* end of ex1.c */
```

The program begins with

```
#include <GL/glut.h>
```

All OpenGL programs must start with this line, which accesses all the OpenGL include files: it pulls in all the function prototypes and other definitions used by OpenGL. Miss it out, and `cogl` will flatly refuse to compile your program.

`ex1.c` contains two functions: `display()`, and `main()`. The execution of all C programs starts at `main()`, so we'll start there too.

We first call the **glutInit()** function:

---

void **glutInit** (int **argc*,
                 char ***argv*);

---

glutInit() initializes the GLUT library, and it must be called before any other GLUT function. `argc` and `argv` should be the arguments of the application's `main()` – **glutInit()** understands several command-line options, which are beyond the scope of this manual (see the GLUT manual for details).

Next, we call **glutCreateWindow()**:

---

int **glutCreateWindow** (char **name*);

---

glutCreateWindow() creates an OpenGL window for rendering and interaction, with `name` displayed in its titlebar. GLUT assigns this window an integer identifier, returned as the result of the function. The window identifier is used when writing OpenGL programs which use multiple windows (described in Section 9.4). By default, the window has a size of $(300, 300)$ pixels, and its position is up to the window manager to choose. If the functions **glutInitWindowSize()** or **glutInitWindowPosition()** (page 19) have already been called, their arguments will control the size and position of the window.

Next comes a call to **glutDisplayFunc()**, and this is a bit more interesting. It's an example of one of the cornerstones of OpenGL programming, which we'll need to look at in detail – the use of **callback functions**.

4.4 Callback functions

A callback function, more often just called a **callback**, is a C function, written by the application programmer. In program `ex1.c`, `display()` is the only callback function we define. But there's one

important difference between a callback function and an ordinary C function: the application never calls the callback function directly. Instead, the callback function is called **by OpenGL**, whenever OpenGL decides it needs to be called.

In `ex1.c`, we use the most basic callback of all – a function that draws the graphics that we want OpenGL to display. We use **glutDisplayFunc()** to tell OpenGL which application function it should call whenever it needs to refresh the window to draw graphics:

---

void **glutDisplayFunc** (*void (*func)(void)*);

---

glutDisplayFunc() registers the name of the callback function to be invoked when OpenGL needs to redisplay (or display for the first time) the contents of the window. The application **must** register a display function – it isn't optional.

The argument of **glutDisplayFunc()** is rather cryptic, and worth a closer look:

```
void (*func)(void)
```

This says that `func()` must be a function which returns `void`, and has no arguments. In other words, a function like `display()`:

```
void display (void) {
/* Called when OpenGL needs to update the display */
  glClear (GL_COLOR_BUFFER_BIT); /* Clear the window */
  glFlush();                     /* Force update of screen */
}
```

So to summarise, in our example the line:

```
glutDisplayFunc (display); /* Register the "display" function */
```

tells OpenGL to call the application's function `display()` function whenever it needs to redraw the graphics.

It's up to the application to define what the `display()` function does – who else could know? In `ex1.c`, the `display()` function doesn't do much: it simply calls **glClear()**:

---

void **glClear** (GLbitfield *mask*);

---

glClear() clears one or more of OpenGL's buffers, specified by `mask`. In this manual, we'll only be concerned with one buffer, the **frame buffer**, which holds the pixels which will be copied to the window. This has the special name GL_COLOR_BUFFER_BIT. When **glClear()** is called, each pixel in the buffer is set to the **current clear colour**, which is set to black by default. You set the current clear colour using the function **glClearColor()** (see page 74).

Now we have a call to **glFlush()**:

---

void **glFlush** (void);

The purpose of this function is to instruct OpenGL to make sure the screen is up to date – it causes the contents of any internal OpenGL buffers are "flushed" to the screen. Note that you only ever

need to call **glFlush()** when you're not using **double-buffering** (which we'll meet in Chapter 6). In practice, most OpenGL programs will use double-buffering – to stop screen flicker – but for now in these simple examples we're not using it just yet.

What would happen if we didn't call **glFlush()** at the end of `display()`? Then, we couldn't guarantee that the screen will show the up-to-date picture. And that's clearly not desirable for a real-time interactive graphics program!

4.5 The main event loop

glutMainLoop() starts the GLUT "event processing" loop:

> void **glutMainLoop** (void);

Once started, this loop will carry on for as long as the program is running. Each time around the loop, GLUT checks to see if anything has changed since last time, and calls the appropriate callback functions.

In pseudocode, the action of **glutMainLoop()** is this:

```
while (1) { /* loop forever */
   if (the application has changed the graphics) {
      call the DISPLAY callback function;
      }

   if (the window has been moved or resized) {
      call the RESHAPE callback function;
      }

   if (any keyboard and/or mouse events have happened) {
      call the KEYBOARD and/or MOUSE callback function;
      }

   call the IDLE callback function;

} /* while */
```

We'll ignore the `reshape()` function for now, returning to it in Section 5.3. And we'll look at the `idle()` function in Section 6.1.

4.6 Example 2: a keyboard event callback

As we saw above, quitting `ex1.c` must be done from the command-line, which isn't very nice from a user-interface point of view. Here's how we can do it better, using a callback, in program `ex2.c`:

```
/* ex2.c */
#include <GL/glut.h>
#include <stdio.h>
```

```
  void display (void) {
  /* Called when OpenGL needs to update the display */
    glClear (GL_COLOR_BUFFER_BIT); /* Clear the window */
    glFlush();                     /* Force update of screen */
  }

  void keyboard (unsigned char key, int x, int y) {
  /* Called when a key is pressed */
    if (key == 27) exit (0);    /* 27 is the Escape key */
    else printf ("You pressed %c\n", key);
  }

  int main(int argc, char **argv) {
    glutInit (&argc, argv);       /* Initialise OpenGL */
    glutCreateWindow ("ex2");     /* Create the window */
    glutDisplayFunc (display);    /* Register the "display" function */
    glutKeyboardFunc (keyboard);  /* Register the "keyboard" function */
    glutMainLoop ();              /* Enter the OpenGL main loop */
    return 0;
  }
  /*end of ex2.c */
```

Try `ex2.c` out.

The addition we've made is to tell OpenGL what to do when it detects a keyboard event. We tell it to call the function `keyboard()` using **glutKeyboardFunc**():

void **glutKeyboardFunc** (*void (*func)(unsigned char key, int x, int y)*);

glutKeyboardFunc() registers the application function to call when OpenGL detects a key press generating an ASCII character. This can only occur when the mouse focus is inside the OpenGL window.

Again, the specification of the argument type is a bit cryptic. It says that it expects a function `func()` which returns `void`, and has the three arguments `key`, `x` and `y`. So, it's a function like this:

```
  void keyboard (unsigned char key, int x, int y) {
  /* Called when a key is pressed */
  }
```

Three values are passed to the callback function: `key` is the ASCII code of the key pressed; `x` and `y` give the pixel position of the mouse at the time.

Back to `ex2.c` – inside the `keyboard()` callback, we look at the value of `key`. If it's 27 (the ASCII code for the `escape` key – surely you knew that!) we call the standard C function `exit()` to terminate the program cleanly; otherwise, we print (in the shell window) a message saying which key was pressed. Note that `ex2.c` needs an extra `#include` line:

```
  #include <stdio.h>
```

because we're using the `printf()` function.

Note: **glutKeyboardFunc()** only responds to pressed keys which have single ASCII codes. For other keys, such as the arrow or function keys, use the **glutSpecialFunc()** function (page 67).

4.7   Example 3: customizing the window

In `ex3.c` we add a few new functions to give us better control over the drawing window:

```
/* ex3.c */
#include <GL/glut.h>

void display (void) {
/* Called when OpenGL needs to update the display */
    glClearColor (1.0,1.0,1.0,0.0);
    glClear (GL_COLOR_BUFFER_BIT); /* Clear the window */
    glFlush();                     /* Force update of screen */
}

void keyboard (unsigned char key, int x, int y) {
/* Called when a key is pressed */
  if (key == 27) exit (0);    /* 27 is the Escape key */
}

int main(int argc, char **argv) {
  glutInit (&argc, argv);       /* Initialise OpenGL */
  glutInitWindowSize (500, 500);    /* Set the window size */
  glutInitWindowPosition (100, 100); /* Set the window position */
  glutCreateWindow ("ex3");     /* Create the window */
  glutDisplayFunc (display);    /* Register the "display" function */
  glutKeyboardFunc (keyboard); /* Register the "keyboard" function */
  glutMainLoop ();              /* Enter the OpenGL main loop */
  return 0;
}
/* end of ex3.c */
```

Try `ex3.c` out.

First, we specify a size and position for the window using **glutInitWindowSize()**:

```
void glutInitWindowSize ( int width,
                          int height );
```

glutInitWindowSize() sets the value of GLUT's **initial window size** to the size specified by `width` and `height`, measured in pixels.

Similarly, **glutInitWindowPosition()** sets the value of GLUT's **initial window position**:

```
void glutInitWindowPosition ( int x,
                              int y );
```

x and y give the position of the top left corner of the window measured in pixels from the **top left corner** of the X display.

4.8   What next?

Now onto Chapter 5, which looks at 2D and 3D graphics.

Chapter 5

2D and 3D graphics

In this chapter we start doing some graphics. We'll begin by extending `ex3.c` to do some 2D drawing – just a triangle, but it'll serve to illustrate how drawing works in OpenGL.

5.1 Example 4: drawing a 2D triangle

`ex4.c` draws a triangle, using the coordinates shown in Figure 5.1.

Figure 5.1: The triangle from example `ex4.c`. It's defined on the $Z = 0$ plane. The Z axis comes out of the page towards you.

Here's the code:

```
/* ex4.c */
#include <GL/glut.h>

void display (void) {
/* Called when OpenGL needs to update the display */
    glClear (GL_COLOR_BUFFER_BIT); /* Clear the window */
    glLoadIdentity ();
    gluLookAt (0.0, 0.0, 0.5, 0.0, 0.0, 0.0, 0.0, 1.0, 0.0);
```

```
      glBegin (GL_LINE_LOOP); /* Draw a triangle */
        glVertex3f(-0.3, -0.3, 0.0);
        glVertex3f(0.0, 0.3, 0.0);
        glVertex3f(0.3, -0.3, 0.0);
      glEnd();
    glFlush();                    /* Force update of screen */
}

void keyboard (unsigned char key, int x, int y) {
/* Called when a key is pressed */
   if (key == 27) exit (0); /* 27 is the Escape key */
}

void reshape (int width, int height)
{ /* Called when the window is created, moved or resized */
    glViewport (0, 0, (GLsizei) width, (GLsizei) height);
    glMatrixMode (GL_PROJECTION); /* Select the projection matrix */
    glLoadIdentity ();             /* Initialise it */
    glOrtho(-1.0,1.0, -1.0,1.0, -1.0,1.0); /* The unit cube */
    glMatrixMode (GL_MODELVIEW);  /* Select the modelview matrix */
}

int main(int argc, char **argv) {
  glutInit (&argc, argv);        /* Initialise OpenGL */
  glutInitWindowSize (500, 500);     /* Set the window size */
  glutInitWindowPosition (100, 100); /* Set the window position */
  glutCreateWindow ("ex4");    /* Create the window */
  glutDisplayFunc (display);   /* Register the "display" function */
  glutReshapeFunc (reshape);   /* Register the "reshape" function */
  glutKeyboardFunc (keyboard); /* Register the "keyboard" function */
  glutMainLoop ();             /* Enter the OpenGL main loop */
  return 0;
}
/* end of ex4.c */
```

Try ex4.c out. You should see a white triangle on a black background.

Although this is a simple example, it illustrates one of the most crucial aspects of OpenGL– **viewing**. OpenGL is a system for drawing 3D graphics. But display screens are 2D – they're flat. Figure 5.2 shows the situation.

In example eg4.c, we draw the triangle on the $Z = 0$ plane. But this is still 3D graphics!

5.2 Viewing using the camera

The idea of creating a 2D view of a 3D scene is simple: we "take a picture" of the scene using a **camera**, and display the camera's picture in the window on the display screen. For convenience, OpenGL splits the process into three separate steps:

- **Step one**: First, we specify the position and orientation of the camera, using the function **glu-**
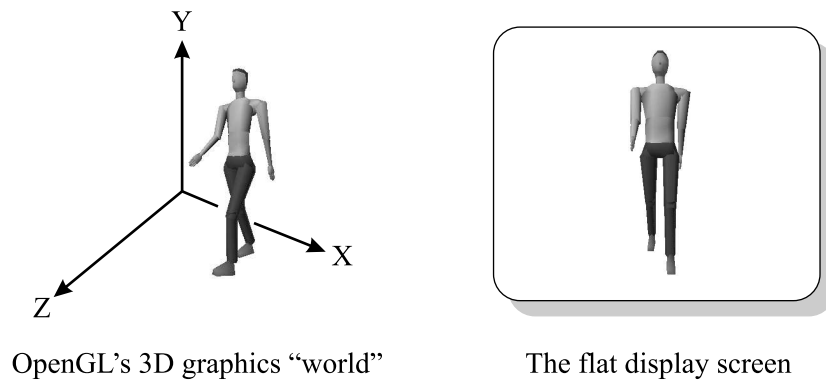
<div align="center">OpenGL's 3D graphics "world"   The flat display screen</div>

<div align="center">Figure 5.2: OpenGL's 3D "world", and the 2D display screen.</div>

LookAt();

- **Step two**: Second, we decide what kind of projection we'd like the camera to create. We can choose an **orthographic** projection (also known as a **parallel projection**) using the function **glOrtho()** (page 56); or a **perspective** projection using the function **gluPerspective()** (page 56);

- **Step three**: Finally, we specify the size and shape of the camera's image we wish to see in the window, using **glViewport()** (page 58). This last step is optional – by default the camera's image is displayed using the whole window.

In OpenGL, the camera model described above is always active – you can't switch it off. It's implemented using **transformation matrices**, and we describe this in detail in Chapter 9. For now, here's a brief description of the process.

OpenGL keeps two transformation matrices: the **modelview** matrix, M; and the **projection matrix**, P. The modelview matrix holds a transformation which composes the scene in world coordinates, and then takes a view of the scene using the camera (step one, above). The projection matrix applies the camera projection (step two, above).

Whenever the application program specifies a coordinate c for drawing, OpenGL transforms the coordinate in two stages, as follows, to give a new coordinate c'. First it transforms the coordinate c by the matrix M, and then by the matrix P, as follows:

$$c' = P \cdot M \cdot c$$

When an OpenGL application starts up, P and M are unit matrices – they apply the **identity transformation** to coordinates, which has no effect on the coordinates. It's **entirely up to the application** to ensure that the M and P matrices always have suitable values. Normally, an application will set M in its `display()` function, and P in its `reshape()` function, as we shall now describe.

5.3 The window reshape function

After creating the window, and registering the display and keyboard callbacks, we now register a new function, the `reshape()` callback:

void **glutReshapeFunc** (*void (*func)(int width, int height)*);

glutReshapeFunc() registers the application callback to call when the window is first created, and also if the window manager subsequently informs OpenGL that the user has reshaped the window. The new height and width of the window, in pixels, are passed to the callback. Typically, the callback will use these values to define the way that OpenGL's virtual camera projects its image onto the window, as we see in the next section.

5.3.1 Specifying the projection

We usually specify the projection in the `reshape()` callback function, because the projection will often need to be adjusted if the user changes the shape of the window. In example `ex4.c` we use an orthographic (also known as "parallel") projection:

```
void reshape (int width, int height)
{ /* Called when the window is created, moved or resized */
  glViewport (0, 0, (GLsizei) width, (GLsizei) height);
  glMatrixMode (GL_PROJECTION); /* Select the projection matrix */
  glLoadIdentity ();
  glOrtho(-1.0,1.0, -1.0,1.0, -1.0,1.0); /* The unit cube */
  glMatrixMode (GL_MODELVIEW);  /* Select the modelview matrix */
}
```

We begin by setting the **viewport** using **glViewport()**, which specifies a rectangular portion of the window in which to display the camera's image. As in this example, it's common to use the the whole of the window, so we set the viewport to be a rectangle of equal dimensions to the window. We'll look at **glViewport()** in detail in Section 9.3.

Next, we set up an orthographic projection. **glMatrixMode()** (page 47) selects which matrix subsequent functions will affect – in this case we select the projection matrix (P). Then we initialise it to the unit transformation with **glLoadIdentity()** (page 48). This is very important, as we shall see in a moment. Then, we select the orthographic projection using **glOrtho()** (page 56). The projection we've chosen maps a unit cube, centred on the origin, onto the viewport.

glOrtho() actually does two things: first it creates a new temporary matrix (let's call it T) to implement the projection, and then it multiplies P with T, as follows:

$$P = P \cdot T$$

That's why we need to make sure P is initialised to the unit transformation first.

Note that the `reshape()` function ends with another call to **glMatrixMode()**, which this time selects the modelview matrix (M) for subsequent modification, for when we position the camera in the `display()` function.

5.3.2 Positioning the camera

This is usually done in the application's `display()` function, using the function **gluLookAt()**. We'll describe this function in detail in Section 9.1. In `ex4.c`, we use it to position the camera on the Z axis at $(0.0, 0.0, 0.5)$, looking towards the origin:

```
glLoadIdentity ();          /* start with a unit modelview matrix */
gluLookAt (0.0, 0.0, 0.5,   /* position of camera */
           0.0, 0.0, 0.0,   /* point at which camera looks */
           0.0, 1.0, 0.0); /* "up" direction of camera */
```

Again, because **gluLookAt()** creates a new transformation and multiplies it into the current matrix (M in this case), we need to ensure that M is first initialised using **glLoadIdentity()**.

5.4 Example 5: a 3D cube with perspective projection

We now turn to 3D drawing, and `ex5.c` draws a cube, centred on the origin:

```
/* ex5.c */
#include <GL/glut.h>

void display (void) {
/* Called when OpenGL needs to update the display */
  glClear (GL_COLOR_BUFFER_BIT); /* Clear the window */
  glLoadIdentity ();
  gluLookAt (0.0, 0.0, 5.0, 0.0, 0.0, 0.0, 0.0, 1.0, 0.0);
  glutWireCube(2.0);
  glFlush();                      /* Force update of screen */
}

void keyboard (unsigned char key, int x, int y) {
/* Called when a key is pressed */
  if (key == 27) exit (0);       /* 27 is the Escape key */
}

void reshape (int w, int h) {
/* Called if the window is moved or resized */
    glViewport (0, 0, (GLsizei)w, (GLsizei)h);
    glMatrixMode (GL_PROJECTION);
    glLoadIdentity ();
    gluPerspective (60, (GLfloat)w / (GLfloat)h, 1.0, 100.0);
    glMatrixMode (GL_MODELVIEW);
}

int main(int argc, char **argv) {
  glutInit (&argc, argv);          /* Initialise OpenGL */
  glutInitWindowSize (500, 500);   /* Set the window size */
  glutInitWindowPosition (100, 100); /* Set the window position */
  glutCreateWindow ("ex5");        /* Create the window */
  glutDisplayFunc (display);       /* Register the "display" function */
```

```
    glutKeyboardFunc (keyboard);        /* Register the "keyboard" function */
    glutReshapeFunc (reshape);          /* Register the "reshape" function */
    glutMainLoop ();                    /* Enter the OpenGL main loop */
    return 0;
}
/* end of ex5.c */
```

Try `ex5.c` out.

In `display()`, we call **glutWireCube()**, which draws a wire-frame cube (see page 42). This time, however, we view it using a **perspective** projection as specified in our `reshape()` function:

```
gluPerspective (60, /* field of view in degrees */
                (GLfloat)w / (GLfloat)h, /* aspect ratio of view */
                1.0, 100.0); /* near and far clipping planes */
```

gluPerspective() sets a perspective projection, so we see the kind of view a camera would normally give, where lines further away from the viewer appear smaller. Here, we specify a field of view of 60 degrees, and an aspect (width-to-height) ratio for the view which exactly matches the aspect ratio of the window. We'll explain the use of clipping planes in Chapter 9.

5.5   What next?

Now onto Chapter 6, which looks at the use of **double buffering** for achieving smooth animation.

Chapter 6

Animated graphics

Computer graphics really comes to life when we draw images that **move**.

6.1   Example 6: a rotating cube

In this next example – `ex6.c` – we'll make OpenGL spin the cube about its centre. Have a look at the code, then take a copy of the program, and compile and run it:

```
/* ex6.c */
#include <GL/glut.h>

GLfloat angle= 0.0;

void spin (void) {
   angle+= 1.0;
   glutPostRedisplay();
}

void display(void) {
   glClear (GL_COLOR_BUFFER_BIT);
   glLoadIdentity ();
   gluLookAt (0.0, 0.0, 5.0, 0.0, 0.0, 0.0, 0.0, 1.0, 0.0);
   glRotatef(angle, 1, 0, 0);
   glRotatef(angle, 0, 1, 0);
   glRotatef(angle, 0, 0, 1);
   glutWireCube(2.0);
   glFlush();               /* Force update of screen */
}

void reshape (int w, int h) {
   glViewport (0, 0, (GLsizei)w, (GLsizei)h);
   glMatrixMode (GL_PROJECTION);
   glLoadIdentity ();
   gluPerspective (60, (GLfloat) w / (GLfloat) h, 1.0, 100.0);
   glMatrixMode (GL_MODELVIEW);
}
```

```
void keyboard(unsigned char key, int x, int y) {
  if (key == 27)  exit (0);                    /* escape key */
}

int main(int argc, char **argv) {
   glutInit(&argc, argv);
   glutInitWindowSize (500, 500);
   glutInitWindowPosition (100, 100);
   glutCreateWindow ("ex6: A rotating cube.");
   glutDisplayFunc(display);
   glutReshapeFunc(reshape);
   glutKeyboardFunc(keyboard);
   glutIdleFunc(spin);            /* Register the "idle" function */
   glutMainLoop();
   return 0;
}
/* end of ex6.c */
```

You should see the cube rotating, but in a rather broken-up sort of way. We'll come back to that in a moment.

The engine behind the animation is the event loop. Using **glutIdleFunc()**, we register an application callback function that gets called each time around the **glutMainLoop()**:

---
void **glutIdleFunc** (*void (*func)(void)*);

---

glutIdleFunc() registers a callback which will be automatically be called by OpenGL in **each cycle** of the event loop, **after** OpenGL has checked for any events and called the relevant callbacks.

In ex6.c, the idle function we've registered is called spin():

```
void spin (void) {
   angle+= 1.0;
   glutPostRedisplay();
}
```

First spin() increments the global variable angle. Then, it calls **glutPostRedisplay()**, which tells OpenGL that the window needs redrawing:

---
void **glutPostRedisplay** (void);

---

glutPostRedisplay() tells OpenGL that the application is asking for the display to be refreshed. OpenGL will call the application's display() callback at the next opportunity, which will be during the next cycle of the event loop.

Note: While OpenGL is processing a single cycle of the event loop, several callbacks may call **glutPostRedisplay()**. Nevertheless, OpenGL won't actually call the display callback until all outstanding events have been dealt with. And, within one cycle of the event loop, a succession of outstanding calls to **glutPostRedisplay()** will be treated as a single call to **glutPostRedisplay()**, so display callbacks will only be executed once – which is probably what you want.
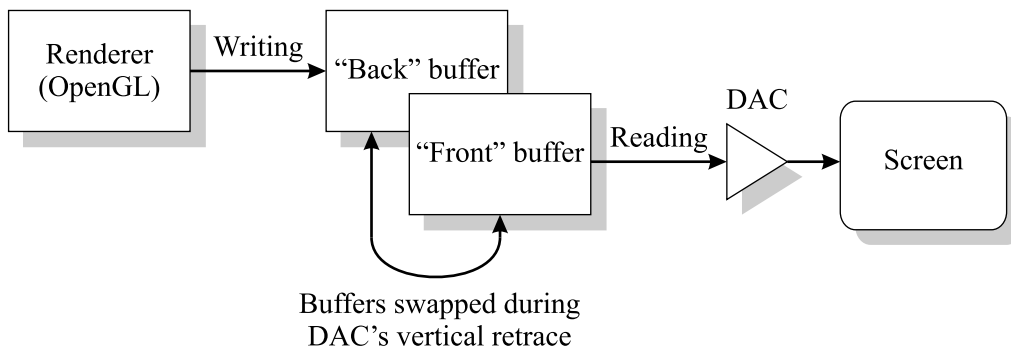
Figure 6.1: Single buffering.

Figure 6.2: Double buffering.

6.2 Double-buffering and animation

As we saw, the rotating cube looks horrible. Why?

The problem is that OpenGL is operating **asynchronously** with the refreshing of the display. OpenGL is pumping out frames too fast: it's writing (into the frame-buffer) a new image of the cube in a slightly rotated position, **before** the previous image has been completely displayed.

Recall the architecture of raster displays: as shown in Figure 6.1, the pixel data is stored in the frame buffer, which is repeatedly read (typically at 60 Hz) by the digital-to-analogue converter (DAC) to control the intensity of the electron beam as it sweeps across the screen, one scan-line at a time. With a single frame-buffer, the renderer (OpenGL) is writing pixel information into the buffer **at the same time** the DAC is reading the information out. If the writer and the reader are out of sync, the reader can never be guaranteed to read and display a complete frame – so the viewer always sees images which comprise part of one frame and part of another. This is very disturbing to the eye – and destroys any possibility of seeing smooth animation.

One solution is to use an additional buffer, as shown in Figure 6.2. The idea here is that one buffer, called the "back buffer" is only ever **written to** by the renderer. The other buffer – the "front buffer" – is only ever **read by** the DAC. The renderer writes its new frame into the back buffer, and when that's done, it then requests that the back and front buffers be swapped over. The **trick** is to perform the swapping while the DAC is performing its **vertical retrace**, which is when it's finished a complete sweep of its buffer, and is resetting to begin again. There's enough slack time here to swap the contents of the two buffers over. This method will ensure that the DAC only ever reads and displays a complete frame.

By default, OpenGL works in **single-buffer** mode, and so we get the fragmented animation seen in

ex6.c. But we can tell OpenGL to use double-buffering, using **glutInitDisplayMode()**:

---
void **glutInitDisplayMode** (unsigned int *mode*);
---

glutInitDisplayMode() sets the **current display mode**, which which will be used for a window created using **glutCreateWindow()**. mode is:

- GLUT_SINGLE: selects a single-buffered window – which is the default if **glutInitDisplayMode** isn't called;

- GLUT_DOUBLE: selects a double-buffered window;

(There are more display modes, beyond the scope of this manual. For a full description, see the GLUT manual or the Red Book.)

For example, to select a double-buffered window you would call:

```
glutInitDisplayMode (GLUT_DOUBLE);
glutCreateWindow ("my window");
```

Once we're using double-buffering, we can tell OpenGL that a frame is complete, and that the buffers should be swapped using **glutSwapBuffers()**:

---
void **glutSwapBuffers** (void);
---

glutSwapBuffers() swaps the back buffer with the front buffer, at the next opportunity, which is normally the next vertical retrace of the monitor. The contents of the new back buffer (which was the old front buffer) are undefined.

Note: Swapping the buffers doesn't have the side effect of **clearing** any buffers. Clearing a buffer must be done explicitly by the application, by calling **glClear()**. Note again that now we're using double-buffering, it's no longer necessary to use **glFlush()**.

6.3   Exercise: smooth the cube

Edit your copy of ex6.c as follows:

- In main(), after the call to **glutInit()**, insert a call to **glutInitDisplayMode()** to select a double-buffered window;

- In display(), after the call to **glutWireCube()**, insert a call to **glutSwapBuffers()**.

- Also, **remove** the call to **glFlush()**. We don't need that anymore, since it gets called internally by **glutSwapBuffers()**. And if we leave **glFlush()** in the code, not only will its effect be redundant, but it'll also slow the program down.

See the difference? Smooth animation!

6.4 Example 7: rotating objects following the mouse

Finally, we now extend `ex6.c` to display a few different objects, and to follow the mouse around.

We won't describe the code here – have a look at `ex7.c` on-line for yourself. And try running it. Cycle between the various objects by pressing the `space` key.

The main new functions we use are **glutPassiveMotionFunc()** (page 68) and **gluUnProject()** (page 59).

6.5 What next?

This is the end of the Tutorial section of the manual. The remaining chapters form the OpenGL Reference Manual.

Part II

OpenGL Reference Manual

Chapter 7

Graphics primitives

In this chapter we describe the coordinate system OpenGL uses, and some of the OpenGL graphics primitives.

7.1   Coordinate systems

OpenGL uses right-handed Cartesian coordinate systems, as shown in Figure 7.1.

Figure 7.1: A right-handed coordinate system. The positive Z axis comes out of the page.

By convention, we draw the positive X axis heading rightwards, the positive Y axis heading vertically, with the positive Z axis heading out of the page towards you.

All the OpenGL functions which create graphical primitives such as lines and polygons work in **object coordinates**. OpenGL automatically transforms object coordinates, first by the **modelview matrix** (M) and then by the **projection matrix** (P). We describe the modelview matrix and the projection matrix in Chapters 8 and 9.

7.2   Defining a vertex

The basic building block for creating graphics with OpenGL is a **point in 3D space**. To describe a shape, you specify the set of points that together make up the shape. In OpenGL terminology, a point in 3D space is called a **vertex**.

You define a single vertex using the function **glVertex3f()**:

```
void glVertex3f ( GLfloat x,
                  GLfloat y,
                  GLfloat z );
```

Here, the "3f" part of the function name means that the function takes three arguments, each of which is a `GLfloat`. As we described in Section 2.6, GL uses its own data types. `GLfloat` is equivalent to the C type `float`.

So, for example, to define the vertex at (10,8,5), shown in Figure 7.1, you would call:

```
glVertex3f (10.0, 8.0, 5.0);
```

7.3   OpenGL function flavours

Many OpenGL functions come in several flavours. For example, suppose you only ever want to do 2D drawing, so you're only concerned with specifying vertices in the XY plane, and all vertices will have a Z coordinate of 0. To make life easier, OpenGL offers a variant form of the **glVertex3f()** function, called **glVertex2f()**:

```
void glVertex2f ( GLfloat x,
                  GLfloat y );
```

Internally, this function still creates a 3D vertex, but it sets its Z coordinate to 0.0 for you, to save you the bother. But in this manual, we will always use the 3D form of functions – the less functions we have to remember, the better!

7.4   Defining shapes: primitives

A vertex on its own isn't very interesting. Now we look at how to group vertices together into **vertex lists**, which define geometrical shapes. The grouping of vertices is done with the **glBegin()** and **glEnd()** functions:

```
void glBegin ( GLenum mode );
```

glBegin() defines the start of a vertex list. `mode` determines the kind of shape the vertices describe, which can be:

- A set of unconnected points (`GL_POINTS`);

- Lines (GL_LINES, GL_LINE_STRIP, GL_LINE_LOOP);

- The boundary of a single convex polygon (GL_POLYGON);

- A collection of triangles (GL_TRIANGLES, GL_TRIANGLE_STRIP, GL_TRIANGLE_FAN);

- A collection of quadrilaterals (GL_QUADS, GL_QUAD_STRIP).

---

void **glEnd** (void);

---

glEnd() defines the end of a vertex list.

7.5 Drawing points

We use the following mode in **glBegin()** to draw points:

- GL_POINTS: each vertex represents a point.

```
glBegin (GL_POINTS);
   glVertex3f (0.0, 6.0, 4.0);
   glVertex3f (0.0, 8.0, 0.0);
   glVertex3f (8.0, 6.0, 0.0);
   glVertex3f (8.0, 3.0, 0.0);
   glVertex3f (6.0, 0.0, 5.0);
   glVertex3f (2.0, 0.0, 5.0);
glEnd ();
```

7.6 Drawing lines

In the function **glBegin()**, the values of mode which interpret vertices as points to connect with lines are:

- GL_LINES: each pair of vertices is drawn as a separate line.

- GL_LINE_STRIP: all the vertices are joined up with lines.

- GL_LINE_LOOP: all the vertices are joined up with lines, and an extra line is drawn from the last vertex to the first.

Figure 7.2 illustrates how the same set of vertices can be drawn as lines in different ways according to mode:

```
glBegin (GL_LINES);  /* or GL_LINE_STRIP or GL_LINE_LOOP */
   glVertex3f (0.0, 6.0, 4.0);
   glVertex3f (0.0, 8.0, 0.0);
   glVertex3f (8.0, 6.0, 0.0);
```

```
    glVertex3f (8.0, 3.0, 0.0);
    glVertex3f (6.0, 0.0, 5.0);
    glVertex3f (2.0, 0.0, 5.0);
glEnd ();
```

Figure 7.2: The same set of vertices drawn using different line styles.

As well as geometry, primitives also have **attributes**, which control ther visual style.

7.6.1   Line attributes

void **glLineWidth** (GLfloat *width*);

glLineWidth() sets the curent line width, measured in pixels. The default value is 1.0.

void **glLineStipple** (GLint *factor*,
                         GLushort *pattern*);

glLineStipple() sets the stippling pattern for lines, which enables lines to be drawn in a flexible variety of dot/dash patterns. By default, stippling is switched off (see Section 7.6.2), and must be enabled by calling:

```
    glEnable(GL_LINE_STIPPLE);
```

Line stippling works on a pixel-by-pixel basis, as the line is rendered into the frame buffer. `pattern` is a 16-bit series of 0s and 1s. When OpenGL renders a line, for each pixel it is about to write, it first consults the next bit in `pattern`, starting at the **low-order bit**, If this bit is a 1, the pixel is written, in the current drawing colour. If the bit is a 0, the pixel is not written.

For example, suppose the pattern specified was (to choose a random example) 0x3E1F. In binary this is:

```
    0011 1110 0001 1111
```

So, when drawing a line, OpenGL would draw the first 5 pixels on, the next 4 off, then one on, the next five on, and the next 2 off. For the next pixel, OpenGL would return to the low-order bit of the pattern, and repeat.

`factor` is a way of elongating the pattern – it multiplies each sub-sequence of consecutive 0s and 1s. For example, if factor=3, then if the bit series 0110 appeared in the pattern, it would be "stretched" to be 01111110.

Handy values of `pattern`, with `factor` set to 1.0, are:

```
Pattern      Rough idea of what the line looks like

0x1111       .    .    .    .    .    .    .

0x3333       ..   ..   ..   ..   ..   ..   ..

0x0F0F       ....     ....      ....      ....

0xAAAA       . . . . . . . . . . . . . .

0xFFFF       ...........................
```

7.6.2 Enabling OpenGL capabilities

void **glEnable** (GLenum *capability*);

void **glDisable** (GLenum *capability*);

OpenGL has a number of capabilities which by default are not active – for reasons of efficiency. These include lighting, texturing, hidden surface removal and line stippling. To use one of these capabilities, it must be explicitly "enabled" by the application, using **glEnable()**. The capability may be subsequently disabled using **glDisable()**. Some of the valid values of `capability` are:

- `GL_LINE_STIPPLE`
- `GL_LIGHTING`
- `GL_FOG`
- `GL_DEPTH_TEST`

7.7 Drawing triangles

The different values of `mode` in **glBegin()** to create triangles are:

- `GL_TRIANGLES`: each triplet of points is drawn as a separate triangle. If the number of vertices is not an exact multiple of 3, the final one or two vertices are ignored.

- `GL_TRIANGLE_STRIP`: constructs a set of triangles with the vertices v0, v1, v2 then v2, v1, v3 then v2, v3, v4 and so on. The ordering is to ensure that the triangles are all drawn correctly form part of surface.

- GL_TRIANGLE_FAN: draws a set of triangles with the vertices v0, v1, v2 then v0, v2, v3 then v0, v3, v4 and so on.

```
glBegin (GL_TRIANGLES);
   glVertex3f (0.0, 6.0, 4.0);
   glVertex3f (0.0, 8.0, 0.0);
   glVertex3f (8.0, 6.0, 0.0);
   glVertex3f (8.0, 3.0, 0.0);
   glVertex3f (6.0, 0.0, 5.0);
   glVertex3f (2.0, 0.0, 5.0);
glEnd ();
```

7.8   Drawing quadrilaterals

We can use two values for mode in **glBegin()** to create quadrilaterals.

- GL_QUADS: each set of four vertices is drawn as a separate quadrilaterals. If the number of vertices is not an exact multiple of 4, the final one, two or three vertices are ignored.

- GL_QUAD_STRIP: constructs a set of quadrilaterals with the vertices v0, v1, v3, v2 then v2, v3, v5, v4 then v4, v5, v7, v6 and so on.

```
glBegin (GL_QUADS); /* or GL_QUAD_STRIP */
   glVertex3f (0.0, 6.0, 4.0);
   glVertex3f (0.0, 8.0, 0.0);
   glVertex3f (8.0, 6.0, 0.0);
   glVertex3f (8.0, 3.0, 0.0);
   glVertex3f (6.0, 0.0, 5.0);
   glVertex3f (2.0, 0.0, 5.0);
glEnd ();
```

7.9   Drawing polygons

We draw a polygon using the following mode in **glBegin()**:

- GL_POLYGON: the vertices define the boundary of a single convex polygon.

The polygon specified must not intersect itself and must be convex. Figure 7.3 shows a polygon with 5 vertices, drawn with the following code:

```
glBegin (GL_POLYGON)
   glVertex3f (0.0, 6.0, 0.0);
   glVertex3f (0.0, 6.0, 6.0);
   glVertex3f (6.0, 6.0, 6.0);
   glVertex3f (9.0, 6.0, 2.0);
   glVertex3f (9.0, 6.0, 0.0);
glEnd ();
```

Figure 7.3: A simple polygon with 5 vertices.

For efficiency and simplicity, OpenGL only guarantees to draw a polygon **correctly** if it's **convex**. A polygon is convex if, taking any pair of points inside the polygon and drawing a straight line between them, all points along the line are also inside the polygon. Figure 7.4 shows a few examples of convex polygons (on the left) and non-convex polygons (on the right).

Note: to draw a non-convex polgyon in OpenGL, it must first be broken into a set of convex polygons, each of which is then drawn separately. This process is called **tesselation**, and non-convex polygons can be broken down this way. GLU provides a set of functions for doing this – see the Red Book, Chapter 11.

Polygons must also be **planar** (completely flat) if they are to be rendered correctly.

Figure 7.4: Convex polygons (left) and non-convex polygons (right).

7.9.1 Polygon attributes

```
void glPolygonMode ( GLenum face,
                     GLenum mode );
```

glPolygonMode() sets the drawing mode for polygons.

face can be GL_FRONT, GL_BACK or GL_FRONT_AND_BACK. mode can be GL_FILL, or GL_LINE.

7.10 GLUT's primitives

GLUT provides a number of functions for easily drawing more complicated objects. Each comes in two versions: **wire** and **solid**. The wire forms are drawn using lines (GL_LINE or GL_LINE_LOOP); the solid forms use polygons (with surface normals, suitable for creating lit, shaded images). Note that these objects do not use display lists (see Chapter 14).

7.10.1 Cube

void **glutWireCube** (GLdouble *size*);

glutWireCube() draws a cube, with edge length size, centred on $(0, 0, 0)$ in object coordinates. Solid version: **glutSolidCube()**.

7.10.2 Sphere

void **glutWireSphere** (GLdouble *radius*,
                         GLint *slices*,
                         GLint *stacks*);

glutWireSphere() draws a sphere, of radius radius, centred on $(0, 0, 0)$ in object coordinates. slices is the number of subdivisions around the Z axis (like lines of longitude); stacks is the number of subdivisions along the Z axis (like lines of latitude). Solid version: **glutSolidSphere()**.

7.10.3 Cone

void **glutWireCone** (GLdouble *base*,
                       GLdouble *height*,
                       GLint *slices*,
                       GLint *stacks*);

glutWireCone() draws a cone, with base radius radius, and height height. The cone is oriented along the Z axis, with the base placed at $Z = 0$, and the apex at $Z =$ height. slices is the number of subdivisions around the Z axis; stacks is the number of subdivisions along the Z axis. Solid

version: **glutSolidCone()**.

---

void **glutWireTorus** (GLdouble *innerRadius*,
                          GLdouble *outerRadius*,
                          GLint *nsides*,
                          GLint *rings*);

---

glutWireTorus() draws a torus centred on $(0, 0, 0)$ in object coordinates. The axis of the torus is aligned with the Z axis. `innerRadius` and `outerRadius` give the inner and outer radii of the torus respectively; `nsides` is the number of sides in each radial section, and `rings` is the number of radial sections. Solid version: **glutSolidTorus()**.

7.10.4 Platonic solids

---

void **glutWireTetrahedron** (void);

---

glutWireTetrahedron() draws a tetrahedron (4-sided regular object) of radius $\sqrt{3}$ centred on $(0, 0, 0)$ in object coordinates. Solid version: **glutSolidTetrahedron()**.

---

void **glutWireOctahedron** (void);

---

glutWireOctahedron() draws an octahedron (8-sided regular object) of radius 1 centred on $(0, 0, 0)$ in object coordinates. Solid version: **glutSolidOctahedron()**.

---

void **glutWireDodecahedron** (void);

---

glutWireDodecahedron() draws a dodecahedron (12-sided regular object) of radius $\sqrt{3}$ centred on $(0, 0, 0)$ in object coordinates. Solid version: **glutSolidDodecahedron()**.

---

void **glutWireIcosahedron** (void);

---

glutWireIcosahedron() draws an icosahedron (20-sided regular object) of radius 1 centred on $(0, 0, 0)$ in object coordinates. Solid version: **glutSolidIcosahedron()**.

7.10.5 Teapot

---

void **glutWireTeapot** (GLdouble *scale*);

---

glutWireTeapot() draws a teapot, scaled by `scale`. Solid version: **glutSolidTeapot()**.

Chapter 8

Modelling using transformations

This chapter is about modelling: we explain how to use transformations to assemble 3D scenes. Chapter 9 explains how to create a view of the scene using the camera model.

8.1 Vectors and matrices

We saw in Section 7.1 that a 3D vertex – a point in space – is represented as x, y, z, mathematically, we write a 3D point as a **column vector**: If we have a point p, we write it as:

$$p \leftarrow \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

You'll notice the extra '1' at the bottom of the vector – this is known as a **homogeneous** representation. To cut a long story short, the use of such vector representations is a mathematical trick which allows all common transformation types to be expressed in a consistent manner using 4×4 matrices.

Warning: beware that some Computer Graphics textbooks represent coordinates as **row vectors**. Using row vectors doesn't change the basic methods used for matrix transformations, but the order in which matrices appear, and their rows and columns, are reversed. Trying to think in terms of both column and row vectors is a recipe for disaster. Stick to column vectors always.

In OpenGL all coordinate transformations are specified using 4×4 matrices. If we transform a point p with a matrix M (for example, a scale by sx, sy, sz), we get a transformed point p', as follows:

$$\begin{bmatrix} x' \\ y' \\ z' \\ 1 \end{bmatrix} \leftarrow \begin{bmatrix} sx & 0 & 0 & 0 \\ 0 & sy & 0 & 0 \\ 0 & 0 & sz & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

or, more succinctly:

$$p' \leftarrow M \cdot p$$

If we subsequently transform p' by another matrix N, to give p'', we have:

$$p'' \leftarrow N \cdot p'$$

so expressing the entire transformation we have:

$$p'' \leftarrow N \cdot M \cdot p$$

8.2   A note about matrix ordering

Notice that the order in which the matrices are written, reading from left to right, is the **reverse** of the order in which their transformations are applied. In the above example, the first transformation applied to p is M, and the transformed point is then subsequently transformed by N.

In general, matrix multiplication is **not commutative**. So, with two matrices M and N,

$$M \cdot N \neq N \cdot M$$

In other words, the **order** in which transformation matrices are applied is crucial.

One of the most common problems in computer graphics, **blank screen syndrome (BSS)**, is often due to incorrectly ordered matrix transformations. Your image has been lovingly computed, but it is being displayed several miles to the West of your display screen; or that tiny blob in the left-hand corner of your screen is your image, compressed into a few pixels.

8.3   Selecting the current matrix

OpenGL maintains two separate 4×4 transformation matrices:

- the **modelview** matrix; this is used to accumulate modelling transformations, and also to specify the position and orientation of the camera;

- the **projection** matrix; this is used to specify how the 3D OpenGL world is transformed into a 2D camera image for subsequent display. The projection matrix performs either a perspective or othographic (parallel) projection.

At any time, one or the other of these matrices is selected for use, and is called the **current matrix**, or sometimes C for short. Most of the OpenGL functions for managing transformations affect the contents of C.

8.4   Setting the current matrix

If our first transformation M represents a scale by (sx, sy, sz), and the second transformation N a translation by (tx, ty, tz), we would code this in OpenGL as follows:
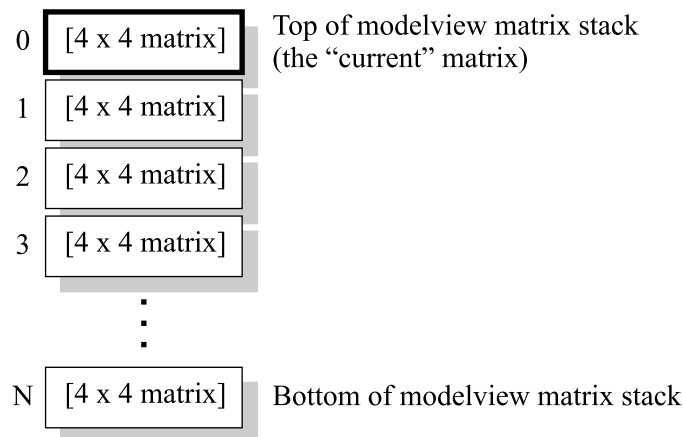
Figure 8.1: An OpenGL modelview matrix stack. The top element of the **selected** stack is often referred to as the "current matrix" C.

```
glMatrixMode (GL_MODELVIEW);  /* Select the modelview matrix */
glLoadIdentity ();            /* Set the current matrix to identity */
glTranslatef (tx, ty, tz);    /* Post-multiply by (tx,ty,tz) */
glScalef (sx, sy, sz);        /* Post-multiply by (sx,sy,sz) */
glVertex3f(x, y, z);          /* Define the vertex */
```

Note that all the OpenGL functions which affect the current matrix C do so by **post-multiplication**. This means that we write the sequence of OpenGL transformation functions in the **reverse order** to the effect they actually have on vertices. This can take a bit of getting used to.

In fact, the modelview matrix isn't a single matrix stored somewhere inside OpenGL – it's actually the top matrix on a **stack** of modelview matrices. This is shown in Figure 8.1. Similarly, the projection matrix is the top matrix on a **stack** of projection matrices. We'll see later why OpenGL uses stacks of matrices.

Figure 8.2 shows how the modelview and projection matrices on the top of their respective stacks affect a vertex specified by the application.

---

void **glMatrixMode** (GLenum *mode*);

---

glMatrixMode() selects the matrix stack, and makes the top matrix on the stack the "current matrix" (C). mode selects the matrix stack, as follows:

- GL_MODELVIEW: selects the modelview matrix stack;

- GL_PROJECTION: selects the projection matrix stack.

Once a current matrix has been selected using **glMatrixMode()**, all subsequent matrix functions (such as **glRotatef()**, etc.) affect the current matrix. For example, to load a translation by (x, y, z) into the modelview matrix, the code would be:
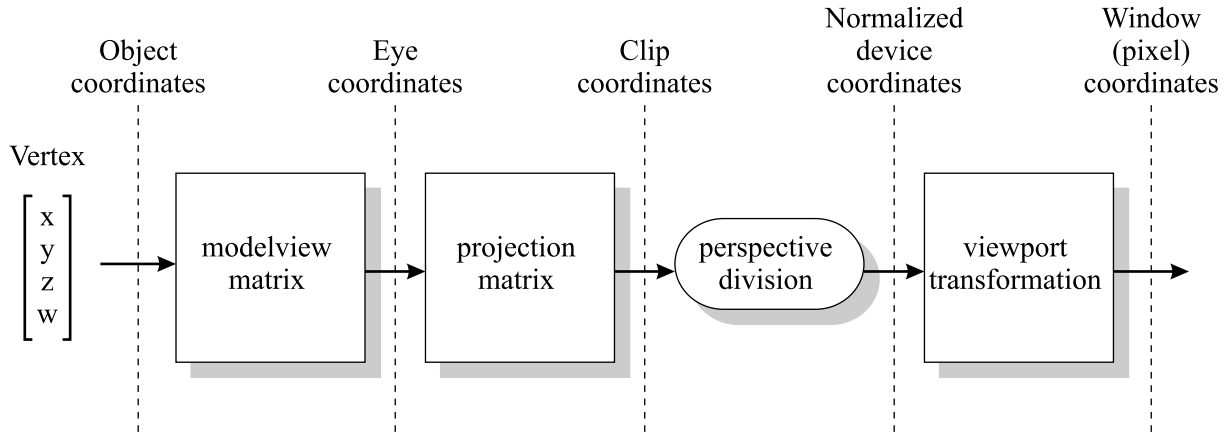
```
glMatrixMode (GL_MODELVIEW);
```

Figure 8.2: The OpenGL viewing pipeline, showing the sequence of transformations and operations applied to a 3D vertex.

```
    glLoadIdentity ();
    glTranslatef (x, y, z);
```

Subsequent matrix operations will continue to affect the current modelview matrix, until **glMatrix-Mode()** is called again to select a different matrix.

8.5   Operating on the current matrix

There are a number of utility functions for changing the current matrix.

8.5.1   Setting to identity

void **glLoadIdentity** (void);

glLoadIdentity() sets the current matrix C to be the identity matrix I:

$$C \leftarrow I$$

where

$$I = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

8.5.2 Translation

```
void glTranslatef ( GLfloat x,
                    GLfloat y,
                    GLfloat z );
```

glTranslatef() creates a matrix M which performs a translation by (x, y, z), and then post-multiplies the current matrix by M as follows:

$$C \leftarrow C \cdot M$$

8.5.3 Scaling

```
void glScalef ( GLfloat x,
                GLfloat y,
                GLfloat z );
```

glScalef() creates a matrix M which performs a scale by (x, y, z), and then post-multiplies the current matrix by M as follows:

$$C \leftarrow C \cdot M$$

8.5.4 Rotation

```
void glRotatef ( GLfloat angle,
                 GLfloat x,
                 GLfloat y,
                 GLfloat z );
```

glRotatef creates a matrix M which performs a counter-clockwise rotation of `angle` degrees. The axis about which the rotation occurs is the vector from the origin $(0, 0, 0)$ to the point (x, y, z), and then post-multiplies the current matrix by M as follows:

$$C \leftarrow C \cdot M$$

8.6 Using the matrix stacks

Because all the OpenGL transformation functions (like **glTranslate()**) always change the current matrix by post-multiplying with the new transformation, sometimes it can be awkward to easily get the correct sequence of transformations. This is where the matrix stacks come in.

There are two separate matrix stacks: one for the modelview matrix and one for the projection matrix. Only one matrix stack is current at a particular time, and this is selected by calling **glMatrixMode()**.

There are two functions which operate on the current matrix stack: **glPushMatrix()** and **glPopMatrix()**. They behave as you might expect:

---
void **glPushMatrix** (void);
---

glPushMatrix() Pushes the current matrix stack down one level. The matrix on the top of the stack is copied into the next-to-top position, as shown in Figure 8.3. The current matrix stack is determined by the most recent call to **glMatrixMode()**. C is not changed. It is an error if **glPushMatrix()** is called when the stack is full.
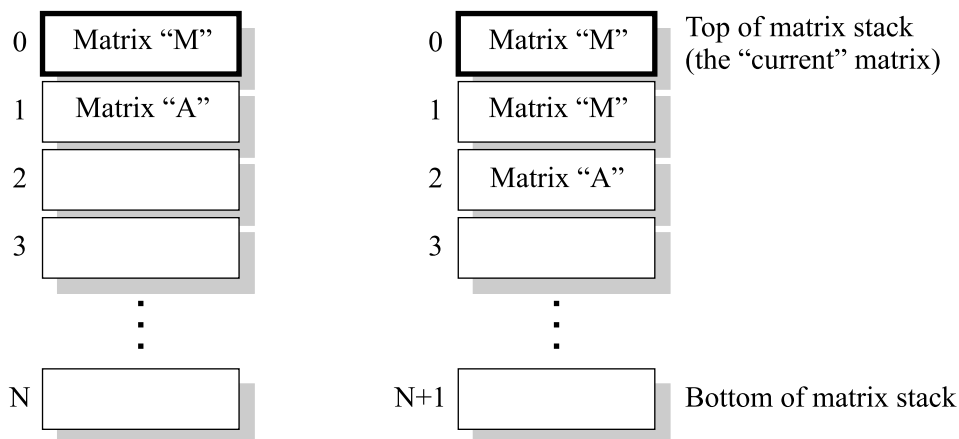
Figure 8.3: The effect of calling **glPushMatrix()** on the current OpenGL matrix stack. The figure shows the stack before (left) and after (right) **glPushMatrix()** is called

Correspondingly, there's a function to pop a matrix off the stack:

---
void **glPopMatrix** (void);
---

glPopMatrix() pops the current matrix stack, moving each matrix in the stack one position towards the top of the stack, as shown in Figure 8.4. The current matrix stack is determined by the most recent call to **glMatrixMode()**. C becomes the matrix previously at the second-to-top of the stack. It is an error if **glPopMatrix()** is called when the stack contains only one matrix.

8.7   Creating arbitrary matrices

You can usually create the matrices you need by using the simple matrix manipulation functions **glLoadIdentity()**, **glTranslate()**, **glScale()** and **glRotate()**, but sometimes – and this is an advanced topic – you need to provide arbitrary 4×4 matrices of your own. See Appendix C for details of how to do this.

0	Matrix "M"		0	Matrix "A"	Top of matrix stack (the "current" matrix)
1	Matrix "A"		1	Matrix "B"	
2	Matrix "B"		2		
3			3		
N			N-1		Bottom of matrix stack
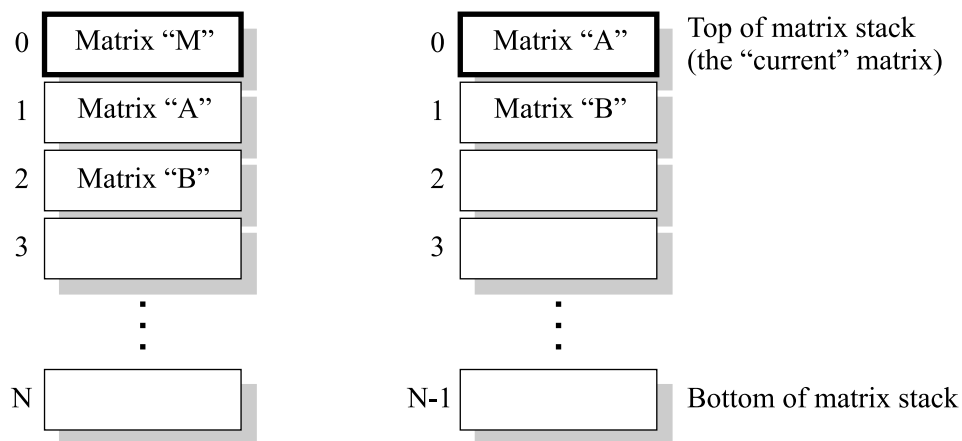
Figure 8.4: The effect of calling **glPopMatrix()** on an OpenGL matrix stack. The figure shows the stack before (left) and after (right) **glPopMatrix()** is called.

Chapter 9

Viewing

In this chapter we look at how to use the OpenGL **viewing model**. The idea is simple: we create a 3D scene using modelling transformations. We then "take a picture" of the scene using a **camera**, and display the camera's picture on the display screen. For convenience, OpenGL splits the process into three separate parts:

- First, we specify the position and orientation of the camera, using **gluLookAt()**.

- Second, we decide what kind of picture we'd like the camera to create. Usually, for 2D graphics we'll use an orthographic (also known as "parallel") view using **glOrtho()**. For 3D viewing, we'll usually want a perspective view, using **gluPerspective()**.

- Finally, we describe how to map the camera's image onto the display screen, using **glViewport()**.

9.1   Controlling the camera

Let's look again at the OpenGL viewing pipeline, in Figure 9.1.

We set the position and orientation of the OpenGL camera, as shown in Figure 9.2, using **gluLookAt()**:

```
void gluLookAt ( GLdouble eyex,
                 GLdouble eyey,
                 GLdouble eyez,
                 GLdouble centerx,
                 GLdouble centery,
                 GLdouble centerz,
                 GLdouble upx,
                 GLdouble upy,
                 GLdouble upz );
```

The position of the camera in space – sometimes also called the **eyepoint** – is given by (eyex, eyey, eyez). (centerx, centery, centerz) specifies a **look point** for the camera to
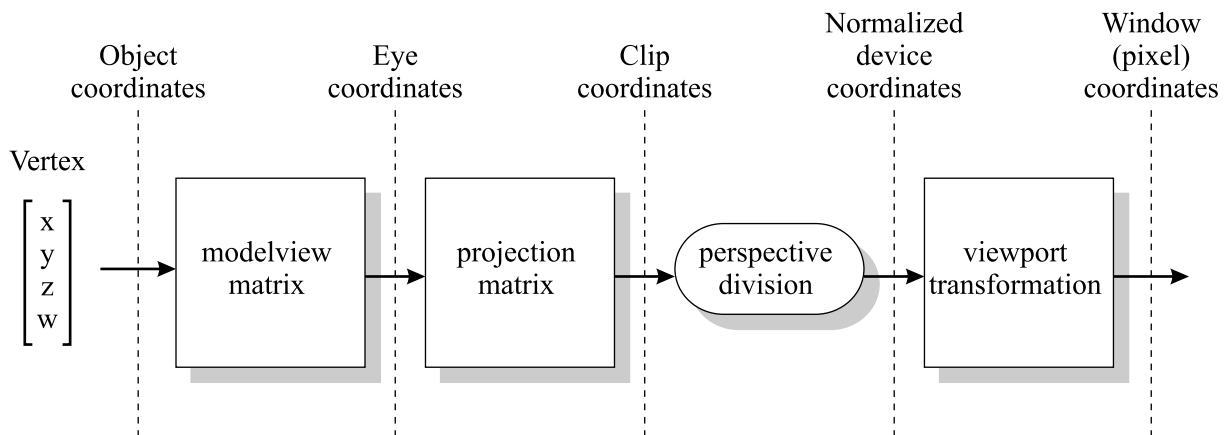
Figure 9.1: The OpenGL viewing pipeline, showing the sequence of transformations and operations applied to a 3D vertex.

"look at", and a good choice for this would a point of interest in the scene, and often the center of the scene is used. Together, the points (`eyex, eyey, eyez`) and (`centerx, centery, centerz`) define a **view vector**. The last set of **gluLookAt()**'s arguments specify the **up vector** of the camera. This defines the camera's orientation at the eyepoint.

There is no need for the view vector and the up vector to be defined at right angles to each other (although if they're parallel weird views may result). Often the up vector is set to a fixed direction in the scene, e.g. pointing up the world Y axis. In the general case, OpenGL twists the camera around the view vector axis until the top of the camera matches the specified up direction as closely as possible.

What **gluLookAt()** actually does is to create a transformation matrix which encapsulates all the specified camera parameters. This is called the "viewing matrix", or V. **gluLookAt()** then post-multiplies the current modelview matrix (C) by V:

$$C \leftarrow C \cdot V$$

If you don't call **gluLookAt()**, the OpenGL camera is given some default settings:

- it's located at the origin, $(0, 0, 0)$;

- it looks down the negative Z axis;

- its "up" direction is parallel to the Y axis.

This is the same as if you had called **gluLookAt()** as follows:

```
gluLookAt (0.0, 0.0, 0.0,   /* camera position */
           0.0, 0.0, -1.0,  /* point of interest */
           0.0, 1.0, 0.0);  /* up direction */
```

Note that the value -1.0 could be any negative float, because this is just specifying the direction "down the negative Z axis".

Figure 9.2: The OpenGL camera.

9.2 Projections

Now that we've positioned and pointed the OpenGL camera, the next step is to specify what kind of image we want. This is done using the **projection matrix**, P. OpenGL applies the projection transformation **after** it has applied the modelview transformation.

9.2.1 The view volume

Consider the real world camera analogy, in which we choose the lens type (wide-angle, telephoto etc.). The choice of lens affects the field of view, and selects what portion of the 3D world will appear within the bounds of final image. The volume of space which eventually appears in the image is known as the **view volume** (or **view frustum**). As well as discarding objects which lie outside the image "frame" OpenGL also imposes limits on how far away objects must be from the camera in order to appear in the final picture.

The actual 3D shape of the view volume depends on what kind of projection is used. For orthographic (parallel) projections the view volume is box-shaped, whereas perspective projections have a view volume shaped like a truncated pyramid. The facets encasing the view volume effectively define six **clipping planes**, which partition the frustum interior from the unseen outside world.

9.2.2   Orthographic projection

glOrtho() creates a matrix for an orthographic projection, and post-multiplies the current matrix (which is normally the projection matrix) by it:

```
void glOrtho ( GLdouble left,
               GLdouble right,
               GLdouble bottom,
               GLdouble top,
               GLdouble near,
               GLdouble far );
```

Figure 9.3 illustrates how the arguments are interpreted. The values define a box-shaped view volume. It is important to set the values such that `left` < `right`, `bottom` < `top` and `near` < `far`. The contents of the view volume are projected onto a rectangular region in the XY plane, with an aspect ratio `(right - left)/(top - bottom)`.

Figure 9.3: The orthographic viewing voulme specified by **glOrtho**.

9.2.3   Perspective projection

gluPerspective() creates a matrix for a perspctive projection, and post-multiplies the current matrix (which will normally be the projection matrix) by it:

```
void gluPerspective ( GLdouble fovy,
                      GLdouble aspect,
                      GLdoublea near,
                      GLdouble far );
```

Figure 9.4 illustrates how the arguments are interpreted. `fovy` is the angle (in degrees) of the image's

vertical field of view; `aspect` is the aspect ratio of the frustum – its width divided by its height; and `near` and `far` respectively specify the positions of the near and far clipping planes, measured as their distances from the **centre of projection** (eyepoint). `near` and `far` must have positive values.

Figure 9.4: The perspective viewing frustum specified by **gluPerspective**.

9.3 Setting the viewport

glViewport() sets the position and size of the **viewport** – the rectangular area in the display window in which the final image is drawn, as shown in Figure 9.5:
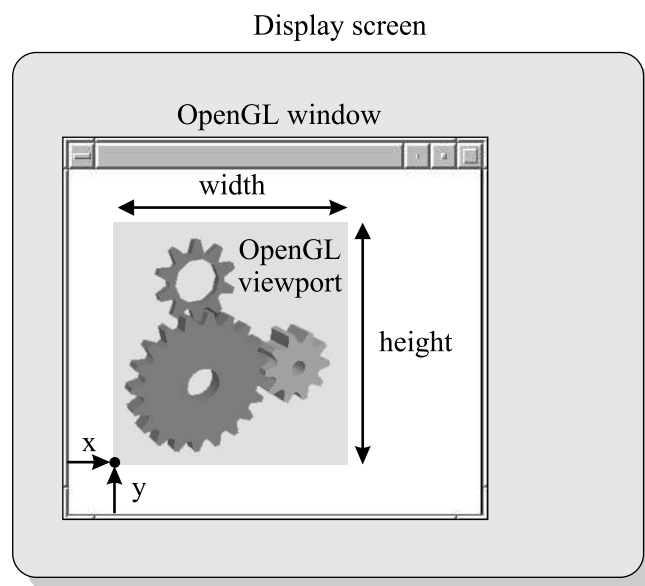
Figure 9.5: How the viewport is defined.

> void **glViewport** (GLint *x,*
> GLint *y,*
> GLsizei *width,*
> GLsizei *height*);

`x` and `y` specify the lower-left corner of the viewport, and `width` and `height` specify its width and height. If a viewport is not set explicitly it defaults to fill the entire OpenGL window. This means that if the window's aspect ratio does not match that defined in **gluPersepctive()** or **glOrtho()** (e.g. after a window resize) the displayed image will appear distorted. **glViewport()** may also be used to draw several separate images within a single OpenGL window.

9.4   Using multiple windows

Most OpenGL programs use a single drawing window. However GLUT does support the use of multiple windows simultaneously. During execution of an OpenGL program, all rendering appears on the **current window**. By default, the current window is always the most recently created window (by **glutCreateWindow()**). If you want to use multiple windows, first create each window and note the window identifier returned by each call to **glutCreateWindow()**. Then, select a window to render using **glutSetWindow()**:

> void **glutSetWindow** (int *window*);

To find out which window is currently selected, call **glutSetWindow()**:

> int **glutGetWindow** (void);

You can also destroy windows, using:

> void **glutDestroyWindow** (int *window*);

Obviously, you can't refer to the window identifier for a window which has been destroyed.

9.5   Reversing the viewing pipeline

Sometimes you'll want to click a pixel point in the window and find out what point in your original object coordinates it corresponds to. This is easy to work out – all you have to do is to invert the viewport, projection and modelview transformations as follows:

Given an object coordinate P_o, its corresponding pixel coordinate P_p is given by:

$$P_p = M_{\text{viewport}} \cdot M_{\text{projection}} \cdot M_{\text{modelview}} \cdot P_o$$

So, if we know P_p, we can obtain P_o by applying the **inverse** of each of the transformations:

$$P_o = M^{-1}_{\text{modelview}} \cdot M^{-1}_{\text{projection}} \cdot M^{-1}_{\text{viewport}} \cdot P_p$$

But there's a problem with doing this. Because a screen pixel position is 2D, and our original object coordinates were 3D, every point along a vector in object coordinates can project to the same 2D screen position. This means it isn't possible to perform an **unambiguous** reverse projection from screen to world. So, the application must choose a z value for the pixel too, which lies between the near and far clipping planes.

---

int **gluUnProject** (GLdouble *winx*,
                         GLdouble *winy*,
                         GLdouble *winz*,
                         const GLdouble *modelMatrix [16]*,
                         const GLdouble *projMatrix [16]*,
                         const GLint *viewport [4]*,
                         GLdouble **objx*,
                         GLdouble **objy*,
                         GLdouble **objz*);

---

gluUnProject() maps the window coordinates `winx`, `winy`, `winz` into object coordinates `objx`, `objy`, `objz`.

The following code, taken from the example program `unproject.c`, shows how **gluUnproject()** is typically used:

```
GLdouble projmatrix[16], mvmatrix[16];
GLint viewport[4];

glGetIntegerv (GL_VIEWPORT, viewport);
glGetDoublev (GL_MODELVIEW_MATRIX, mvmatrix);
glGetDoublev (GL_PROJECTION_MATRIX, projmatrix);
/*  note viewport[3] is height of window in pixels  */
realy = viewport[3] - (GLint) y - 1;
printf ("Coordinates at cursor are (%4d, %4d)\n", x, realy);
gluUnProject ((GLdouble) x, (GLdouble) realy, 0.0,
             mvmatrix, projmatrix, viewport, &wx, &wy, &wz);
printf ("World coords at z=0.0 are (%f, %f, %f)\n",
        wx, wy, wz);
gluUnProject ((GLdouble) x, (GLdouble) realy, 1.0,
             mvmatrix, projmatrix, viewport, &wx, &wy, &wz);
printf ("World coords at z=1.0 are (%f, %f, %f)\n",
        wx, wy, wz);
```

See Section 15.1 for descriptions of the functions **glGetIntegerv()** and **glGetDoublev()**.

Chapter 10

Drawing pixels and images

Sometimes, for image processing applications, you need access pixels directly. Often the most convenient way to do this is to set up a view which gives a one-to-one mapping between object coordinates and pixel coordinates.

10.1   Using object coordinates as pixel coordinates

To do this, you define an orthographic projection, where the width and height of the viewing volume **exactly match** the width and height of the viewport. Normally you will draw on the $z = 0$ plane, so we set the near and far clipping planes to -1.0 and 1.0 respectively.

To begin with, let's assume your program's `main` has created an OpenGLwindow of an appropriate size:

```
glutInitWindowSize (360, 335);
glutInitWindowPosition (100, 100);
glutCreateWindow ("Pixel world");
```

It's usual to place the projection specification in the `reshape` function:

```
void reshape (int width, int height)
{
    glViewport (0, 0, (GLsizei) width, (GLsizei) height);
    glMatrixMode (GL_PROJECTION);
    glLoadIdentity ();
    glOrtho (0.0, (GLfloat) width, 0.0, (GLfloat) height, -1.0, 1.0);
    glMatrixMode (GL_MODELVIEW);
    glLoadIdentity ();
}
```

then the object coordinate point $(60.0, 40.0, 0.0)$ would map to the OpenGLwindow pixel at $(60, 40)$.
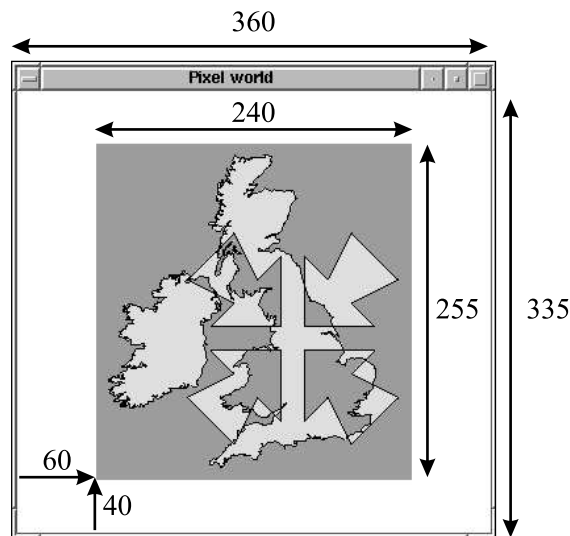
Figure 10.1: The pixel rectangle drawn by **glDrawPixels** at the current raster position.

10.2   Setting the pixel drawing position

The function **glRasterPos3f()** sets the **current raster position** – the pixel position at which the next pixel rectangle specified using **glDrawPixels()** will be drawn:

```
void glRasterPos3f ( GLfloat x,
                     GLfloat y,
                     GLfloat z );
```

The position (x, y, z) is expressed in object coordinates, and is transformed in the normal way by the modelview and projection matrices.

10.3   Drawing pixels

glDrawPixels draws a rectangle of pixels, with `width` pixels horizontally, and `height` pixels vertically.

```
void glDrawPixels ( GLsizei width,
                    GLsizei height,
                    GLenum format,
                    GLenum type,
                    const GLvoid *pixels );
```

The bottom left-hand corner of the pixel rectangle is positioned at the **current raster position**.

`pixels` is a pointer to an array containing the actual pixel data. Because pixel data can be encoded in several different ways, the type of `pixels` is a (`void *`) pointer. `format` and `type` specify the

pixel data encoding: normally `format` will be `GL_RGB`, which states that each pixel is described by three sequential values giving the red, green and blue componets; `type` specifies the data type used for each of the R, G and B components, and will normally be `GL_FLOAT`, with each of the R, G and B values in the range [0.0,1.0].

For example, used in conjunction with the viewing code in Section 10.1, the following code defines and draws a pixel rectangle of size 240 by 255, positioned at (60,40). The result is shown in Figure 10.1:

```
#define WIDTH  240
#define HEIGHT 255

GLfloat image[WIDTH][HEIGHT][3]; /* pixel data, R,G,B */

/* code omitted to write pixel values into 'image' */

void display (void)
{
  glClear(GL_COLOR_BUFFER_BIT);
  glRasterPos3f(60.0, 40.0, 0.0);
  glDrawPixels(WIDTH, HEIGHT, GL_RGB, GL_FLOAT, image);
}
```

Chapter 11

Displaying text

Unlike many graphics systems, OpenGL doesn't directly support the specification and rendering of text. It's left up to the application programmer to draw text using one of two approaches:

- Define the shape of a character as collection of pixels in a bitmap. Here, the shape of a character is not geometric, and so it isn't affected at all by the modelview and projection matrices.

- Draw the shape of each character using OpenGL primitives, most commonly lines. With this approach, each character is a little geometrical object, and can be transformed using the modelview and projection matrices like ordinary OpenGL primitives.

Clearly, both of these methods mean that the application programmer would have to do quite a lot of work to draw text! Fortunately, however, the GLUT library comes to the rescue.

GLUT provides a number of font definitions in both the bitmap and line (also known as "stroke") forms. Here, we only describe GLUT's bitmap text, which is most commonly used. For details of GLUT's stroke text, see the GLUT manual.

11.1   GLUT's bitmap fonts

GLUT defines three groups of bitmap fonts, based on standard X-windows fonts:

- A fixed-width font, where each character occupies a pixel rectangle of fixed size, either 9×15 or 8×13: `GLUT_BITMAP_9_BY_15`, `GLUT_BITMAP_8_BY_13`;

- A proportionally-spaced Times-Roman font, at 10 or 24 points: `GLUT_BITMAP_TIMES_ROMAN_10`, `GLUT_BITMAP_TIMES_ROMAN_24`;

- A proportionally-spaced Helvetica font, at 10, 12 or 24 points: `GLUT_BITMAP_HELVETICA_10`, `GLUT_BITMAP_HELVETICA_12`, `GLUT_BITMAP_HELVETICA_18`.

You can see these fonts (and a stroke font) demonstrated in the example program `font.c`. Use the arrow keys to rotate the fonts – you'll see that the stroke font is properly transformed, but the only the start points of the bitmap fonts move.

11.2   Drawing a single character

void **glutBitmapCharacter** (void **font*,
                                          int *char*);

glutBitmapCharacter() draws the single character whose ASCII code is `char`, from font `font`. The position at which the character's bitmap is drawn is the **current raster position**, set by **glRasterPos3f**().

11.3   Drawing a text string

An application will often wish to display text strings. Here's a simple function to do that (taken from `thegears.c`):

```
void drawString (void *font, float x, float y, char *str) {
/* Draws string 'str' in font 'font', at world (x,y,0) */
   char *ch;
   glRasterPos3f(x, y, 0.0);
   for (ch= str; *ch; ch++)
       glutBitmapCharacter(font, (int)*ch);
}
```

We might call this function as follows, to draw a string at world $(-7.0, 0.0, 0.0)$:

```
glColor3f(1.0, 1.0, 1.0); /* Select white */
drawString (GLUT_BITMAP_HELVETICA_18, -7.0, 0.0, "Press Esc to quit");
```

Note: make sure the Z-coordinate of your bitmap text remains inside the frustum. If it doesn't, OpenGL will clip the entire text out. (Note also that if you're doing lighting (see Chapter 16) – you'll have to ensure lighting is disabled while you're drawing bitmap text.)

Chapter 12

Interaction

The basic OpenGL library has no facilities for interaction – it's only concerned with rendering. This was a design decision made in the interests of efficiency and portability.

The GLUT library provides some very rudimentary facilities for creating graphical user interfaces (GUIs). Specifically, the **glutMainLoop()** function traps events, and allows an application to deal with them in three ways:

- **Mouse events** are triggered when a mouse button is pressed, and also when the mouse changes position;

- **Keyboard events** are triggered when the user hits an ASCII key or a cursor movement/function key;

- **Menu events** are triggered when the application has defined GLUT pop-up menus and assigned them to mouse buttons.

12.1   Keyboard events

For keyboard events, GLUT calls the application callback function registered by **glutKeyboardFunc()** or **glutSpecialFunc()**:

> void **glutKeyboardFunc** (*void (*func)(unsigned char key, int x, int y)*);

glutKeyboardFunc() registers the application function to call when OpenGL detects a key press generating an ASCII character. This can only occur when the mouse focus is inside the OpenGL window.

> void **glutSpecialFunc** (*void (*func)(int key, int x, int y)*);

glutSpecialFunc() registers the application callback to call when OpenGL detects a that key press generating a non-ASCII character has occurred. This can only occur when the mouse focus is inside the OpenGL window. Three values are passed to the callback: `key` is an integer code for the key pressed; `x` and `y` give the pixel position of the mouse. Some useful codes are:

```
GLUT_KEY_LEFT    Left arrow key
GLUT_KEY_RIGHT   Right arrow key
GLUT_KEY_UP      Up arrow key
GLUT_KEY_DOWN    Down arrow key
GLUT_KEY_F1      F1 function key (and similarly F2-F12)
```

12.2   Mouse events

---
void **glutMouseFunc** (*void (*func)(int button, int state, int x, int y)*);

---

glutMouseFunc() registers an application callback function which GLUT will call when the user presses a mouse button within the window. The following values are passed to the callback function:

- button records which button was pressed, and can be

  - GLUT_LEFT_BUTTON
  - GLUT_MIDDLE_BUTTON
  - GLUT_RIGHT_BUTTON

- state records whether the event was generated by pressing the button (GLUT_DOWN), or releasing it (GLUT_UP).

- x, y give the current mouse position in pixels. Note: when using OpenGL with X, the mouse y position is measured from the **top** of the window.

---
void **glutMotionFunc** (*void (*func)(int x, int y)*);

---

glutMotionFunc() registers an application callback function which GLUT will call when the mouse moves within the window while one of its buttons is pressed. The current mouse position x, y is passed to the callback function.

---
void **glutPassiveMotionFunc** (*void (*func)(int x, int y)*);

---

glutPassiveMotionFunc() has the same job as **glutMotionFunc()**, but no buttons need to be pressed for an event to be generated.

12.3   Controlling the mouse cursor

You can set the position of the cursor using **glutWarpPointer()**:

---
void **glutWarpPointer** (int *x*,
                          int *y*);

---

where (x,y) is in pixels relative to the window's origin (top-left).

To change the shape of the mouse cursor, use **glutSetCursor()**:

---

void **glutSetCursor** (int *cursor*);

Where `cursor` is one of the following:

```
GLUT_CURSOR_NONE          /* Turns the cursor off */

GLUT_CURSOR_RIGHT_ARROW /* Basic arrows */
GLUT_CURSOR_LEFT_ARROW

GLUT_CURSOR_INFO          /* Symbolic cursor shapes */
GLUT_CURSOR_DESTROY
GLUT_CURSOR_HELP
GLUT_CURSOR_CYCLE
GLUT_CURSOR_SPRAY
GLUT_CURSOR_WAIT
GLUT_CURSOR_TEXT
GLUT_CURSOR_CROSSHAIR

GLUT_CURSOR_UP_DOWN       /* Directional cursors */
GLUT_CURSOR_LEFT_RIGHT

GLUT_CURSOR_TOP_SIDE     /* Sizing cursors */
GLUT_CURSOR_BOTTOM_SIDE
GLUT_CURSOR_LEFT_SIDE
GLUT_CURSOR_RIGHT_SIDE
GLUT_CURSOR_TOP_LEFT_CORNER
GLUT_CURSOR_TOP_RIGHT_CORNER
GLUT_CURSOR_BOTTOM_RIGHT_CORNER
GLUT_CURSOR_BOTTOM_LEFT_CORNER

GLUT_CURSOR_INHERIT      /* Inherit from parent window */

GLUT_CURSOR_FULL_CROSSHAIR /* Fullscreen crosshair (if available) */
```

12.4 Menu events

GLUT menus are very straightforward to use. Once defined by the application and attached to a specified mouse button, they pop up on the window at the position of the mouse when the appropriate button is pressed. The user then makes a selection from the items in the menu, and GLUT calls an application callback function for the menu, passing as an argument the number of the selected item. GLUT menus can have items which invoke pop-up sub-menus.

The following example program, `menu.c` shows how to create two menus, one attached to the right mouse button, and one to the middle mouse button. Try running the program.

```
/* menu.c */
#include <GL/glut.h>
#include <stdio.h>
```

```
void display (void)
{ /* Callback called when OpenGL needs to update the display */
  glClear (GL_COLOR_BUFFER_BIT); /* Clear the window */
}

void keyboard (unsigned char key, int x, int y)
{ /* Callback called when a key is pressed */
  if (key == 27) { exit (0); }         /* 27 is the Escape key */
}

void tobys_bistro (int item)
{ /* Callback called when the user clicks the right mouse button */
  printf ("Toby's bistro: you clicked item %d\n", item);
}

void steves_chippy (int item)
{ /* Callback called when the user clicks the middle mouse button */
  printf ("Steve's chippy: you clicked item %d\n", item);
}

int main (int argc, char** argv)
{
  glutInit (&argc, argv);          /* Initialise OpenGL */
  glutCreateWindow ("Menus");      /* Create the window */
  glutDisplayFunc (display);       /* Register the "display" function */
  glutKeyboardFunc (keyboard);     /* Register the "keyboard" function */

  glutCreateMenu (tobys_bistro); /* Create the first menu & add items */
  glutAddMenuEntry ("Petto di Tacchino alla Napoletana", 1);
  glutAddMenuEntry ("Bruschetta al Pomodoro e Olive", 2);
  glutAddMenuEntry ("Chianti Classico", 3);
  glutAttachMenu (GLUT_RIGHT_BUTTON); /* Attach it to the right button */

  glutCreateMenu (steves_chippy); /* Create the second menu & add items */
  glutAddMenuEntry ("Rissoles", 1);
  glutAddMenuEntry ("Curry sauce", 2);
  glutAddMenuEntry ("Vimto", 3);
  glutAttachMenu (GLUT_MIDDLE_BUTTON); /* Attach it to the middle button */

  glutMainLoop ();                       /* Enter the OpenGL main loop */
  return 0;
}
/* end of menu.c */
```

12.4.1   Defining menus

---

int **glutCreateMenu** (*void (*func) (int value)*);

---

glutCreateMenu() creates a new pop-up menu, which becomes the **current menu**. The argument is

the name of the application's callback function which is called when an item in the menu is selected by the user. **glutCreateMenu()** allocates the new menu a unique `int` identifier number, which it returns.

> int **glutAddMenuEntry** (char *name*,
>   int *value*);

glutAddMenuEntry() adds new item to the end of the current menu. `name` is the text to display in the item. `value` is the value passed to the application's callback if this item is selected by the user.

> void **glutAddSubMenu** (char *name*,
>   int *menu*);

glutAddSubMenu() adds a new sub-menu to the end of the current menu. `name` is the text to display in the item in the current menu which, when pressed, will display the sub-menu. `menu` is the identifer of the sub-menu, which is created separately with a call to **glutCreateMenu()**.

> void **glutAttachMenu** (int *button*);

glutAttachMenu() attaches the current menu to mouse button `button`. Whenever this button is subsequently pressed, the menu will pop up. `button` must be one of:

- `GLUT_LEFT_BUTTON`
- `GLUT_MIDDLE_BUTTON`
- `GLUT_RIGHT_BUTTON`

> void **glutSetMenu** (int *menu*);

glutSetMenu() sets the **current menu** to the menu whose identifier is `menu`.

12.4.2 Changing menus dynamically

It's also possible to change the items in a menu as the program runs.

> void **glutChangeToMenuEntry** (int *entry*,
>   char * *name*,
>   int *value*);

glutChangeToMenuEntry() changes an entry in the current menu.

Chapter 13

Colour

We express colour using a **colour model**, which gives us a way of assigning numerical values to colours. A common simple model is the Red-Green-Blue (RGB) model, where a colour is represented by a mixture of the three primary colours red, green and blue. This is illustrated in Figure 13.1.

Figure 13.1: The RGB colour model.

13.1   RGB colour in OpenGL

OpenGL supports the RGB colour model, in a slightly extended form. OpenGL adds a fourth component to the colour, called **alpha**, and the revised model is called the **RGBA** model. Alpha represents the opacity (or, equivalently the transparency) of a colour, and is used when blending colours together.

We don't discuss the use of alpha further in this manual.

> void **glClearColor** (GLclampf *red*,
> GLclampf *green*,
> GLclampf *blue*,
> GLclampf *alpha*);

glClearColor() sets the current clearing colour to be used when clearing a buffer using **glClear()**. red, green and blue are the RGB components of the colour. The GLclampf datatype limits these values to floats in the range $[0.0, 1.0]$. Set alpha to 0.0.

> void **glColor3f** (GLclampf *red*,
> GLclampf *green*,
> GLclampf *blue*);

glColor3f() sets the current drawing colour, using a triple of RGB values in the range $[0.0, 1.0]$.

Chapter 14

Retained data

Graphics systems typically act in two ways:

- **Immediate mode**: whenever an application **defines** a primitive, it is **drawn** immediately.

- **Retained mode**: the actions of **defining** a primitive and **drawing** a primitive are treated quite separately. When an application defines a primitive, the graphics system keeps a record of the definition as a **display list**, but the primitive isn't drawn. Subsequently, the application requests that the stored primitive be drawn.

OpenGL provides both ways of working.

14.1   Immediate mode vs retained mode

By default, OpenGL works in **immediate mode**: whenever a primitive is defined, OpenGL draws it immediately. Once it has been drawn and rendered as pixels, OpenGL forgets all about the original primitive.

For example, if we execute the code:

```
glBegin(GL_TRIANGLES)
    glVertex(1.0, 3.0, 0.0);
    glVertex(5.0, 3.0, 0.0);
    glVertex(3.0, 4.0, 0.0);
glEnd();
```

OpenGL will draw the triangle defined by the three vertices, once they have been transformed by the graphics pipeline, as pixels in the display buffer. But OpenGL does not keep any internal record of the original definition of the vertices in object coordinates.

The use of immediate mode has a very important consequence for the application programmer: to ensure that the contents of display are up-to-date, the application must execute all the code that defines primitives (including setting transformations and rendering parameters).

There's another way to use OpenGL, called **retained mode**, which is quite different from **immediate mode**.

14.2   Retained mode

Here, the graphical shapes which are to be drawn are specified within a display list. The OpenGL display list mechanism is best thought of as a **cache** for graphics. It isn't a full-fledged data structure which the application can manipulate. Once created, a display list:

- cannot be edited – its data is execute-only.

- cannot be queried – an application cannot "read back" the data stored in a display list.

If display lists sound very restricted in their functionality, that's exactly the intention. The OpenGL display list is designed for efficiency, not versatility.

14.3   Using display lists

glNewList() creates and opens a new display list named `list`:

```
void glNewList ( GLuint list,
                 GLenum mode );
```

All subsequent OpenGL commands will be stored in the display list. The `list` argument is a positive integer which identifies the display list being created. It is up to the programmer to allocate unique `list` values for each display list. `mode` determines what happens while the list is being created. There are two options:

- `GL_COMPILE`: the commands are not executed as they are stored in the display list. This means the contents of the display list will not be drawn until the display list is called using **glCallList()**.

- `GL_COMPILE_AND_EXECUTE`: the commands are executed as soon as they are stored in the display list.

Only one display list can be open for writing at a time: OpenGL will report an error if another display list is already open. If a display list with the name `list` already exists when, **glNewList()** is called, OpenGL automatically empties the existing display list, and overwrites it with the new definition.

```
void glEndList ( void );
```

glEndList() closes the currently open display list, marking the end of its definition.

The following example creates a simple display list which draws a green triangle:

```
GLuint TRI= 1;

glNewList (TRI, GL_COMPILE);
   glBegin (GL_TRIANGLES);
   glColor3f (0.0, 1.0, 0.0);  /* Green */
   glVertex3f (1.0, 3.0, 0.0);
```

```
      glVertex3f (5.0, 3.0, 0.0);
      glVertex3f (3.0, 4.0, 0.0);
      glEnd ();
   glEndList ();
```

Note that here we've used a symbolic constant `TRI` for the name of the display list, rather than the raw number `1`. This helps readability.

Once a display list has been created, it can be instanced repeatedly, using **glCallList**():

---

void **glCallList** (GLuint *list*);

The effect of calling **glCallList()** on a display list named `list` is to execute again all the OpenGL commands stored in the list. Any drawing specified by the commands will happen, as will any changes to the OpenGL context – such as matrices and attributes.

Here's how we could instance the `TRI` display list:

```
   for (i= 0; i < 5; i++) {    /* Instance a display list */
      glTranslatef (0.1, 0.0, 0.0);
      glCallList (TRI);
   }
```

This code will display five instances of `TRI`, each instance shifted in x by 0.1 units.

14.4   Mixing immediate mode with retained mode

It's perfectly acceptable to use immediate mode and retained mode at the same time. Consider the following code fragment:

```
   glColor3f (1.0, 0.0, 0.0);  /* Red */

   glBegin (GL_LINES);          /* Draw a line (immediate mode) */
      glVertex3f (0.0, 0.0, 0.0);
      glVertex3f (0.2, 0.5, 0.0);
   glEnd ();

   for (i= 0; i < 5; i++) {    /* Instance a display list */
      glTranslatef (0.1, 0.0, 0.0);
      glCallList (TRI);
   }

   glBegin (GL_LINES);          /* Draw another line (immediate mode) */
      glVertex3f (0.0, 0.0, 0.0);
      glVertex3f (0.5, 0.5, 0.0);
   glEnd ();
```

Here, we first set the current colour to red, then draw a line in immediate mode. Next, as in the previous example, we draw five instances of the green triangle – and note how the **glColor3f()** call

(green) stored in the `TRI` display list overwrites the effect of the colour currently in effect when the display list is called (red). Finally, we draw another line in immediate mode. **Question:** what colour will the second line be? **Answer:** green, because green was the most recent colour selected (when the `TRI` display list was called).

Chapter 15

State

OpenGL is a **state machine**: calling OpenGL functions change the state of the machine. This is a fancy way of saying that inside the OpenGL system are a bunch of global variables which the application can set and query. The current values of these variables control the way OpenGL behaves.

For example, calling **glColor3f()** sets the **current drawing colour**, which will be used for drawing primitives.

15.1 State enquiries

An application can query the values of state variables using a simple "keyword and value" model. There's a separate enquiry function for each **type** of state variable. For example, for integers:

---
int **glGetIntegerv** (GLenum *pname*,
                    GLint **params*);

---

glGetIntegerv() enquires the integer state variable specified by `pname`. For example:

```
GLint col[1];

glGetIntegerv(GL_CURRENT_COLOR, col);
```

The value of the current drawing colour will be returned in `col`. Note that this is an **array** variable.

Similarly, the function **glGetDoublev()** enquires the current value of a `GLdouble` state variable:

---
int **glGetDoublev** (GLenum *pname*,
                    GLdouble **params*);

There are similar functions for enquiring other types of state variable. And there are a huge number of state variables – see Appendix B of the Red Book for a complete list.

15.2   Enquiring the viewing state

Sometimes it's necessary to enquire the current values of the viewing state. This can be done as follows:

```
GLdouble projection[16], modelview[16];
GLint viewport[4];

glGetIntegerv (GL_VIEWPORT, viewport);
glGetDoublev (GL_MODELVIEW_MATRIX, modelview);
glGetDoublev (GL_PROJECTION_MATRIX, projection);
```

Chapter 16

Lighting

This chapter describes the facilities OpenGL provides for lighting and shading 3D scenes, so that the objects in them look solid and (somewhat) realistic. Our intention here is to provide enough background and details to get you going, creating lit and shaded scenes. As always, for further information refer to the Red Book, in particular Chapter 5.

16.1   The OpenGL lighting model

OpenGL provides a "local" lighting model, which computes the illumination of a single polygon with respect to one or more light sources. It needs the following information in order to do this:

- The position of each light source

- The colour of each light source

- How the intensity of each light source decreases with distance

- The type of each light source: ambient, diffuse or specular

- The geometry of the polygon

- The colour of the polygon

In the real world, the interaction between light and matter is incredibly complicated. Much research has been undertaken into techniques for simulating these interactions using mathematics and computer graphics, and there are several sophisticated techniques which can create very realistic images. Two of the best known are ray tracing and radiosity. These are called "global" models, because they consider not only the interactions of one light and one object, but also the interactions between all lights and all objects in the scene. Such interactions are responsible, for example, for reflections and shadows.

The basic OpenGL lighting model cannot compute reflections and shadows, because it considers each polygon in isolation from all others. This may come as a surprise, since many OpenGL-based games clearly display sophisticated illumination. In many cases they do this by computing their own illumination, and rendering it using OpenGL textures.

Nevertheless, OpenGL's lighting model is useful, robust, and fast.

16.2   Hidden surface removal

If 3D scenes are to look plausible, we need to worry about "hidden surface removal". If we have
a scene containing several objects, and we view the scene from a certain position, one object might
obscure another. In order to display such a situation realistically, we must ensure that obscured parts
of objects do not get drawn. By default, OpenGL simply draws objects in the order specified by the
programmer, taking no account of whether one object would obscure another for a given viewpoint.

OpenGL implements hidden-surface removal using a simple technique called depth-buffering (also
known as Z-buffering). This takes place during rasterization, using a "depth buffer" – an array which
records a depth value corresponding to each pixel in the window. Initially, each depth value is set to
be a very large number. Whenever a new pixel is generated, for example during the scan-conversion
of a polygon P_1, the pixel's Z value is compared with the corresponding value in the depth-buffer. If
the pixel's depth is less than that in the buffer, the pixel is drawn and its depth recorded in the depth
buffer, over-writing the previous value. Otherwise, the pixel is not drawn and the depth buffer is not
updated.

Now, suppose subsequently that during the scan-conversion of polygon P_2, the same pixel is gener-
ated, because P_1 and P_2 overlap in the scene. If the depth value of P_2's pixel is greater than that stored
in the depth buffer, then P_2 is further away from the eye than P_1, and so is obscured by P_1.

To tell OpenGL to perform hidden-surface removal using a depth buffer, you need to do three things.

First, in the call to **glutInitDisplayMode()**, instruct GLUT to create a depth buffer, specifying GLUT_DEPTH
in addition to any other flags you're using:

```
glutInitDisplayMode (GLUT_DOUBLE | GLUT_DEPTH);
```

Second, enable the depth test, which is switched off by default, using **glEnable()**:

```
glEnable (GL_DEPTH_TEST);
```

Finally, you need to explictly clear the depth buffer (in other words, re-load it with large depth values)
each time around the rendering loop:

```
void display () {
  glClear (GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT);
  / * all your display code */
}
```

Here's ex8.c, which draws a tumbling green cube orbiting a stationary red cube.

```
#include <GL/glut.h>

  float r= 0.0;
  int hidden= 0;

  void init(void) {
     glClearColor (0.0, 0.0, 0.0, 0.0);
  }
```

```
void spin (void) {
   r+= 1.0;
   glutPostRedisplay();
}

void display(void) {
   glClear (GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT);

   if (hidden) glEnable(GL_DEPTH_TEST);
   else glDisable(GL_DEPTH_TEST);

   glLoadIdentity ();
   gluLookAt (0.0, 0.0, 5.0, 0.0, 0.0, 0.0, 0.0, 1.0, 0.0);

   glColor3f (1.0, 0.0, 0.0);
   glutSolidCube(1.0);           /* Red cube */

   glRotatef(r*2.0, 0, 1, 0);    /* Orbit angle */
   glTranslatef(0.0, 0.0, 1.0); /* Orbit radius */
   glRotatef(r, 1, 0, 0);        /* Tumble in x,y,z */
   glRotatef(r, 0, 1, 0);
   glRotatef(r, 0, 0, 1);
   glColor3f (0.0, 1.0, 0.0);
   glutSolidCube(0.5);           /* Green cube */

   glutSwapBuffers();
}

void reshape (int w, int h) {
   glViewport (0, 0, (GLsizei) w, (GLsizei) h);
   glMatrixMode (GL_PROJECTION);
   glLoadIdentity ();
   gluPerspective (60, (GLfloat) w / (GLfloat) h, 1.0, 100.0);
   glMatrixMode (GL_MODELVIEW);
}

void keyboard(unsigned char key, int x, int y) {
  if (key == 27) { exit (0); }                  /* escape key */
  if (key == 'h')
     hidden= !hidden;
}

int main(int argc, char **argv) {
   glutInit(&argc, argv);
   glutInitDisplayMode (GLUT_DOUBLE | GLUT_RGB | GLUT_DEPTH);
   glutInitWindowSize (500, 500);
   glutInitWindowPosition (100, 100);
   glutCreateWindow ("ex8: Press 'h' to toggle hidden surface removal.");
   init ();
   glutDisplayFunc (display);
   glutReshapeFunc (reshape);
   glutKeyboardFunc (keyboard);
```

```
   glutIdleFunc (spin);
   glutMainLoop ();
   return 0;
}
/* end of ex8.c */
```

By default, hidden-surface removal is off, so the cubes are drawn in the order they're coded in
`display()` – that's the stationary red cube first, then the rotating green cube. Where they over-
lap, the green cube's pixels will always over-write the red cube's pixels.

Press 'h' to switch on hidden-surface removal, and you can now see when the green cube orbits **behind**
the red cube, and is therefore obscured by it.

Notice that the green cube is drawn in a single colour – and doesn't look at all "solid". We'll see how
to address that in subsequent sections.

16.3   Defining lights

By default, lighting is off. It's enabled as follows:

```
glEnable (GL_LIGHTING);
```

OpenGL provides at least eight lights, named GL_LIGHT0 through GL_LIGHT7. By default, each
light is switched off, so a light must be enabled if it is to have any effect. For example, to use light 0:

```
glEnable (GL_LIGHT0);
```

As well as enabling a light, you need to set its position, colour and other attributes, but if you don't,
the light has handy default values. In particular, its colour is white, and its located at the position
$(0, 0, 1)$. We can use these defaults to add a light to the previous example, which we altered slightly
so that the orbiting object is now a sphere. Here's ex9.c:

```
/* ex9.c */
#include <GL/glut.h>

float r= 0.0;
int hidden= 1, flat= 1;

void init(void) {
   glClearColor (0.0, 0.0, 0.0, 0.0);
   glEnable (GL_LIGHTING);
   glEnable (GL_LIGHT0);
}

void spin (void) {
   r+= 1.0;
   glutPostRedisplay();
}
```

```
void display(void) {
    glClear (GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT);

    if (hidden) glEnable(GL_DEPTH_TEST);
    else glDisable(GL_DEPTH_TEST);

    if (flat) glShadeModel (GL_FLAT);
    else glShadeModel (GL_SMOOTH);

    glLoadIdentity ();
    gluLookAt (0.0, 0.0, 5.0, 0.0, 0.0, 0.0, 0.0, 1.0, 0.0);

    glColor3f (1.0, 0.0, 0.0);
    glutSolidCube(1.0);              /* Red cube */

    glRotatef(r*2.0, 0, 1, 0);    /* Orbit angle */
    glTranslatef(0.0, 0.0, 1.0); /* Orbit radius */
    glRotatef(r, 1, 0, 0);          /* Tumble in x,y,z */
    glRotatef(r, 0, 1, 0);
    glRotatef(r, 0, 0, 1);
    glColor3f (0.0, 1.0, 0.0);
    glutSolidSphere(0.5, 20, 15); /* Green sphere */

    glutSwapBuffers();
}

void reshape (int w, int h) {
    glViewport (0, 0, (GLsizei) w, (GLsizei) h);
    glMatrixMode (GL_PROJECTION);
    glLoadIdentity ();
    gluPerspective (60, (GLfloat) w / (GLfloat) h, 1.0, 100.0);
    glMatrixMode (GL_MODELVIEW);
}

void keyboard(unsigned char key, int x, int y) {
  if (key == 27) { exit (0); } /* escape key */
  if (key == 'h') hidden= !hidden;
  if (key == 's') flat= !flat;
}

int main(int argc, char **argv) {
    glutInit(&argc, argv);
    glutInitDisplayMode (GLUT_DOUBLE | GLUT_RGB | GLUT_DEPTH);
    glutInitWindowSize (500, 500);
    glutInitWindowPosition (100, 100);
    glutCreateWindow ("ex9");
    init ();
    glutDisplayFunc (display);
    glutReshapeFunc (reshape);
    glutKeyboardFunc (keyboard);
    glutIdleFunc (spin);
    glutMainLoop ();
    return 0;
```

```
}
/* end of ex9.c */
```

Running `ex9`, you'll see the orbiting sphere lit by `LIGHT0`, and now its different faces are shaded according to how they're oriented with respect to the light source.

Note that you can still use 'h' to toggle hidden surface removal. Try toggling it and observe the incorrect results when it's switched off. For lighting to work correctly, the depth buffer must be **enabled**.

16.4   Defining the shading model

In `ex9` we make use of a new function, **glShadeModel()**:

void **glShadeModel** (GLenum *mode*);

glShadeModel() specifies how OpenGL renders primitives. For example, when rendering a polygon, if `mode` is `GL_FLAT`, OpenGL chooses one vertex of the polygon, computes its colour, and assigns this colour to all pixels in the polygon. If `mode` is `GL_SMOOTH`, OpenGL computes a colour for each vertex, and the interior pixels of the polygon are coloured by interpolating between the vertex colours. If **glShadeModel()** is not called, the default behaviour is `GL_SMOOTH`.

In `ex9`, we initially select `GL_FLAT`, and the 's' key toggles between this and `GL_SMOOTH`.

16.5   Defining materials

There's another difference between the visual appearances of `ex8` and `ex9`. In `ex8`, we use **glColor3f()** to set the colour of the two objects, so they're drawn red and green. In `ex9`, these colour settings are still there, but now they don't have any effect; instead, the objects take their colour only from the light source. When lighting is enabled, we need to specify the intrinsic colours of objects in a more sophisticated way.

OpenGL uses the concept of **material properties**. An object is considered to have a surface made of a material with several properties, which determine how it interacts with light. These are:

- ambient colour: how well the material reflects ambient light

- diffuse colour: how well the material reflects diffuse light

- specular colour: how well the material reflects specular light

- emissiveness: whether the material emits light itself

- shininess: how glossy the material is

There's a single function for setting material properties:

```
void glMaterialfv ( GLenum face,
                    GLenum paramName,
                    TYPE *param );
```

`face` specifies which face of a primitive the material property should effect, and may be GL␣FRONT, GL␣BACK or GL␣FRONT␣AND␣BACK.

`paramName` selects which material property to change, as follows:

- GL␣AMBIENT sets the ambient colour (default is (0.2, 0.2, 0.2, 1.0));

- GL␣DIFFUSE sets diffuse colour (default is (0.8, 0.8, 0.8, 1.0));

- GL␣SPECULAR sets the specular colour (default is (0.0, 0.0, 0.0, 1.0));

- GL␣EMISSION sets the emissive colour (default (0.0, 0.0, 0.0, 1.0));

- GL␣SHININESS sets the specular exponent (default 0.0);

`param` is the value to set. For GL␣SHININESS, the type of this argument is GLfloat; for all other values of `mode`, it's GLfloat *.

Example program `ex10.c` illustrates the use of **glMaterialfv()**, to set the diffuse colours of the cube and sphere. The following extract shows the relevant code:

```
GLfloat redDiffuseMaterial[] = {1.0, 0.0, 0.0, 0.0};
GLfloat greenDiffuseMaterial[] = {0.0, 1.0, 0.0, 0.0};

/* code omitted */

glMaterialfv(GL_FRONT, GL_DIFFUSE, redDiffuseMaterial);
glutSolidCube(1.0);             /* Red cube */

/* code omitted */

glMaterialfv(GL_FRONT, GL_DIFFUSE, greenDiffuseMaterial);
glutSolidSphere(0.5, 20, 15); /* Green sphere */
```

16.6 Defining lights

The properties of a light are defined in a similar way to those of materials, using **glLightfv()**:

```
void glLightfv ( GLenum light,
                 GLenum paramName,
                 TYPE *param );
```

`paramName` selects which light property to change, as follows:

- `GL_AMBIENT` sets the ambient light colour (default (0.0, 0.0, 0.0, 1.0))

- `GL_DIFFUSE` sets the diffuse light colour (default (1.0, 1.0, 1.0, 1.0))

- `GL_SPECULAR` sets the specular light colour (default (1.0, 1.0, 1.0, 1.0))

- `GL_POSITION` sets the position of the light (default (0.0, 0.0, 1.0, 0.0)). If the w coordinate of the position is 0.0, the light is considered to be at $+\infty$, and the (x, y, z) components of its position give the direction the light shines in.

- `GL_SPOT_DIRECTION` sets the direction of a spotlight (default (0.0, 0.0, -1.0))

- `GL_SPOT_EXPONENT` sets the exponent of a spotlight (default 0.0)

- `GL_SPOT_CUTOFF` sets the spotlight cutoff angle. The default value is 180.0, which indicates the light is not a spotlight. Any other value indicates the light is a spotlight.

- `GL_CONSTANT_ATTENUATION` sets the constant attenuation factor (default 1.0)

- `GL_LINEAR_ATTENUATION` sets the linear attenuation factor (default 0.0)

16.7   The lighting equation

OpenGL uses the following lighting equation to compute the colour V of a vertex:

$$V = M_e + (I_g * M_a) + \sum_{i=0}^{n-1} A_i * S_i * (\text{amb}_i + \text{diff}_i + \text{spec}_i)$$

where:

- M_e is the material's emission

- I_g is the scaled global ambient light

- M_a is the material's ambient reflectivity

- A_i is the attentuation factor for light i: $\frac{1}{k_{c_i} + k_{l_i} d_i + k_{q_i} d_i^2}$

- where k_{c_i} is the value of `GL_CONSTANT_ATTENUATION`

- k_{l_i} is `GL_LINEAR_ATTENUATION`

- k_{q_i} is `GL_QUADRATIC_ATTENUATION`

- d_i is the distance from light source i to the vertex.

- S_i is the spotlight effect of light i. If the light isn't a spotlight, $S_i = 1.0$; if the light is a spotlight, but V is outside the light's cone of illumination, $S_i = 0.0$; otherwise, $S_i = \max(\hat{v}_i \cdot d_i, 0)^{n_s}$, where $\hat{v}$ is the normalised vector from the position of light i to vertex V, d_i is the direction of light i, and n_s is `GL_SPOT_EXPONENT`.

- amb_i is the ambient reflection component: $I_{a_i} * M_a$

- diff_i is the diffuse reflection component: $\max(\hat{L}_i \cdot \hat{N}, 0) * I_{d_i} * M_d$

- spec_i is the specular reflection component: $\max(\hat{S}_i \cdot \hat{N}, 0)^n * I_{s_i} * M_s)$

- M_d is the material's diffuse reflectivity

- M_s is the material's specular reflectivity

- I_{a_i} is the ambient component of light i

- I_{d_i} is the diffuse component of light i

- I_{s_i} is the specular component of light i

- $\hat{L}_i$ is the normalised vector from V to the position of light source i

- $\hat{N}$ is the unit normal vector for V

- $\hat{S}_i$ is the normalised vector sum of $\hat{L}_i$ and the normalised vector pointing from V to the view-point.

Appendix A

The cogl script

`cogl` is handy for compling a single OpenGL program, which is normally sufficient for simple applications.

```perl
#!/usr/bin/perl
#
# This is for compiling and linking C
# programs with Mesa, on Linux.
#
# usage: cogl [-g] file.c
#
# Toby Howard, 5 November 1998, version 2
#
$CC= "gcc";
$CFLAGS= "-O3 -fomit-frame-pointer -m486 -Wall -pipe -DFX -DXMESA ";
#
$LIB_PATHS= "-L/usr/X11/lib -L/usr/X11R6/lib ";
$HDR_PATHS= "-I/usr/X11R6/include ";
$FORM_LIBS= "-lforms ";
$OGL_LIBS= "-lglut -lMesaGL -lMesaGLU ";
$X_LIBS= "-lX11 -lXext -lXmu -lXt -lXi ";
#
$LIBS= $OGL_LIBS . $FORM_LIBS . $X_LIBS . "-lm ";

local $fin, $fout;

@dirs = split(/\//,$0); $O = "$dirs[$#dirs]"; # get the program name into $O.

if (($#ARGV < 0) || ($#ARGV > 1)) { &usage; } # only one or two args

if ($#ARGV == 0) { # one arg
   $fin= $ARGV[0];
   }
elsif ($#ARGV == 1) { # two args
      if ($ARGV[0] ne "-g" ) { &usage; }
      else {
         $CFLAGS .= "-g";
         $fin= $ARGV[1];
```

```
      }
    }

if ($fin =~ m/\.c$/) { # only accept file.c
   $fout= $fin; $fout =~ s/\.c$//; # Duff! Must be a nicer way.
   }
else { &usage; }

print "$O v2, 5/11/98: compiling $fin; output program will be: $fout\n";

local ($ret)= system("$CC $CFLAGS $HDR_PATHS $fin $LIB_PATHS $LIBS -o $fout") >> 8;
if ($ret) {die ("$O: gcc failed.\n");}

sub usage {
   print "usage: cogl [-g] file.c\n";
   exit (1);
} # usage

# that's it.
```

You can find cogl on-line at:

```
   /usr/local/bin/cogl
```

Appendix B

Using a makefile

`cogl` is handy for compling a single OpenGL program, which is normally sufficient for simple applications. For more complex projects, however, which split functions across several files, it's better to use a makefile.

We won't discuss here the general principles of makefiles – that's a whole topic in itself – but here's a sample makefile for accessing the Mesa libraries on the Linux teaching system:

```
INCDIR = /usr/include
LIBDIR = /usr/lib
XLIBS = -L/usr/X11/lib -L/usr/X11R6/lib -lX11 -lXext -lXmu -lXt -lXi
GL_LIBS = -L$(LIBDIR) -lglut -lMesaGLU -lMesaGL -lm $(XLIBS)
CC = gcc
CFLAGS = -I${INCDIR} -O3 -fomit-frame-pointer -m486 -Wall -pipe

gears:  gears.o
        ${CC} ${CFLAGS} gears.o -o gears ${GL_LIBS}
```

You can find this makefile on-line at:

```
/opt/info/courses/OpenGL/Makefile
```

Appendix C

Advanced matrix operations

You can usually create the matrices you want by using the simple matrix manipulation functions **glLoadIdentity**(), **glTranslate**(), **glScale**() and **glRotate**(), but sometimes you need to provide arbitrary 4×4 matrices of your own. The functions described in this section enable you to do this. Refer to Section C.1 for details of how OpenGL interprets the sequence of elements in an arbitrary matrix.

> void **glLoadMatrixf** (const GLfloat *m*);

glLoadMatrixf() takes a matrix m (a pointer to a sequence of 16 floats) and sets the current matrix C to this matrix:

$$C \leftarrow m$$

glMultMatrixf() takes a matrix m (a pointer to a sequence of 16 floats) and post-multiplies it with the current matrix C, as follows:

$$C \leftarrow C \cdot m$$

> void **glMultMatrixf** (const GLfloat *m*);

C.1   How an OpenGL matrix is stored

By using the utility functions such as **glRotatef**(), **glMultMatrixf**(), and so on, it's simple to create and manipulate matrices. Some applications, however, may wish to create their own matrices, and pass them to OpenGL.

In order to do this correctly, it's necessary to know how OpenGL stores its matrices internally.

Suppose you wanted to create your own matrix and pass it to OpenGL. We'll take the simple example

of a matrix to perform a translation by (x, y, z), which has the mathematical form:

$$\begin{bmatrix} 1 & 0 & 0 & x \\ 0 & 1 & 0 & y \\ 0 & 0 & 1 & z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

We would normally declare such a matrix in C as follows:

```
/* assume x, y and z are already declared */

float M[4][4]= { 1, 0, 0, x,
                 0, 1, 0, y,
                 0, 0, 1, z,
                 0, 0, 0, 1 };
```

C stores multi-dimensional arrays in **row-major** format, so M is actually this sequence of 16 floats (decimal points omitted for clarity):

```
{ 1, 0, 0, x, 0, 1, 0, y, 0, 0, 1, z, 0, 0, 0, 1 }
```

OpenGL, however expects matrices to be in **column-major** format, where an ordered sequence of elements e_1 through e_{16} defines the following matrix:

$$\begin{bmatrix} e_1 & e_5 & e_9 & e_{13} \\ e_2 & e_6 & e_{10} & e_{14} \\ e_3 & e_7 & e_{11} & e_{15} \\ e_4 & e_8 & e_{12} & e_{16} \end{bmatrix}$$

This is the transpose of row-major format. So, if we pass the matrix M to OpenGL, as the argument to **glLoadMatrixf()** or **glMultMatrixf()**, we won't get the result we expect. OpenGL would access the 16 elements of M "column-wise" and create the following OpenGL matrix:

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ x & y & z & 1 \end{bmatrix}$$

This is a matrix for three-point perspective – not translation! The results will be spectacular, but spectacularly wrong.

The safest thing to do is to stick to OpenGL's functions for manipulating matrices – then you need never worry about the way they're stored. But if you do need to compute exotic matrices and pass them to OpenGL, be very careful with the row/column ordering.

Index

The names of OpenGL functions are printed in **bold**, and a bold page number indicates the main description of the function.